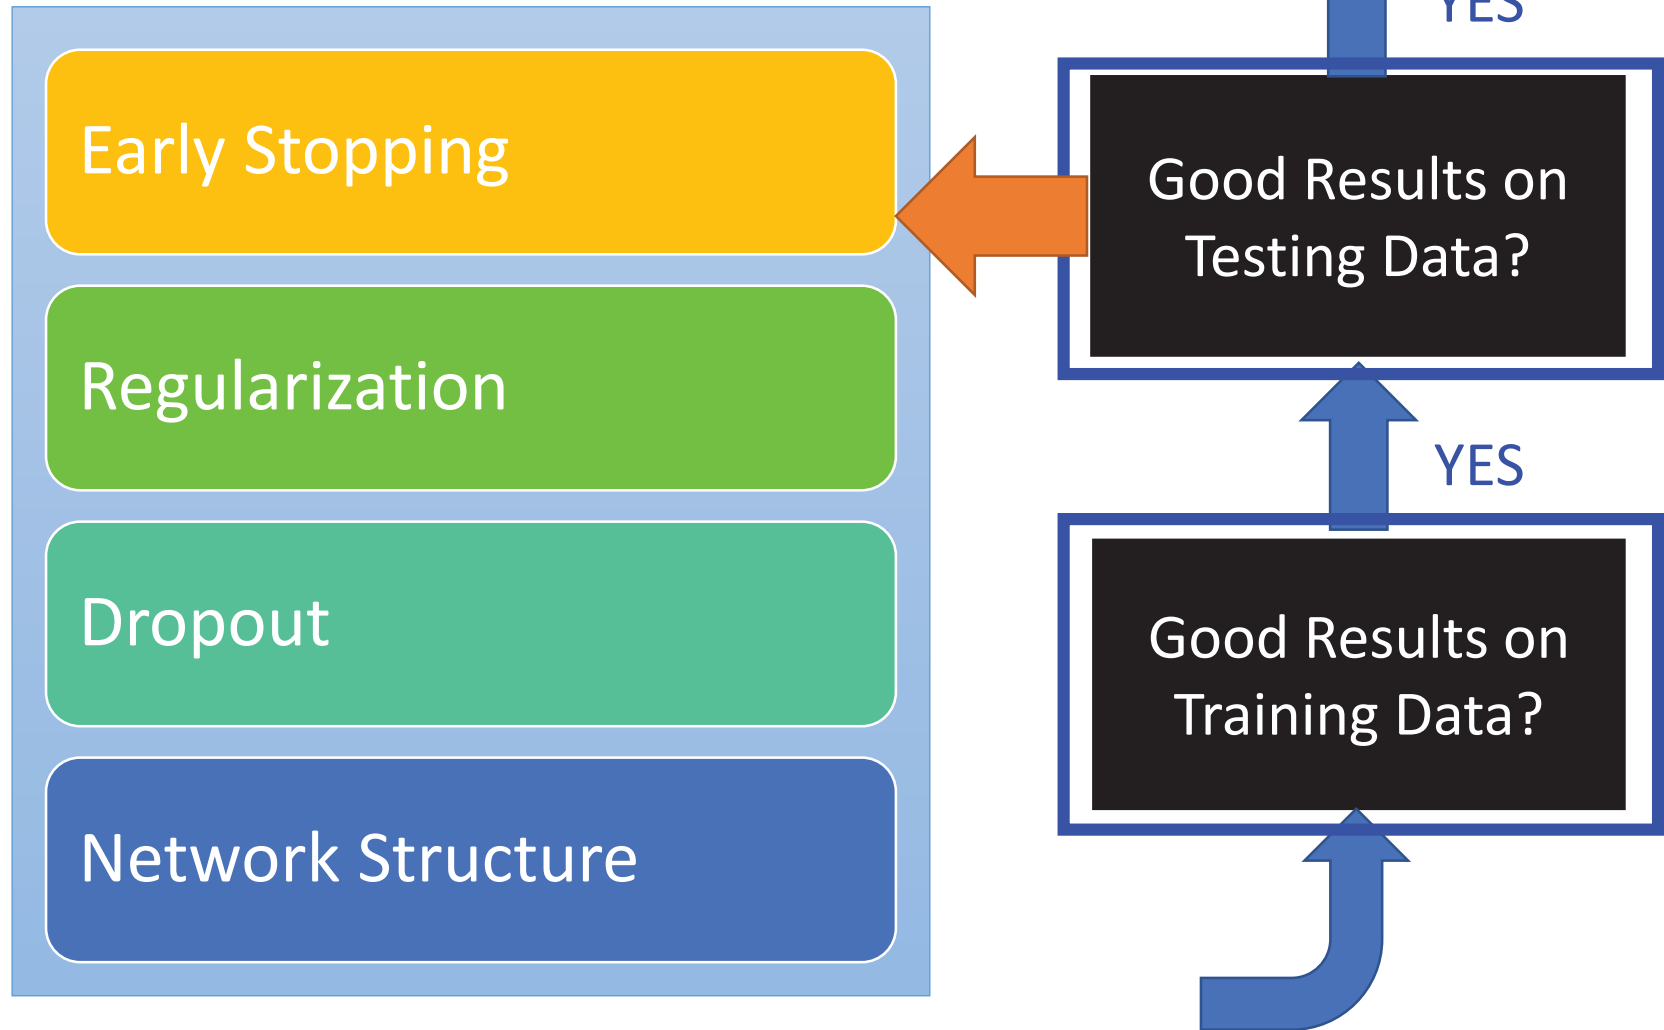



# *Recipe of Deep Learning*




# Why Overfitting?

- Training data and testing data can be different.

Training Data:



Testing Data:



Learning target is defined by the training data.


The parameters achieving the learning target do not necessary have good results on the testing data.

# Panacea for Overfitting


- Have more training data
- **Create** more training data (?)

Handwriting recognition:

Original  
Training Data:



Created  
Training Data:



Shift 15 °

# Why Overfitting?

- For experiments, we added some noises to the testing data

```
-1.36230370e-01, 1.03749340e-01, 1.15432226e-01,  
2.58670464e-01, 1.48774333e+00, 1.92885328e+00,  
1.70038673e+00, 2.46242981e+00, 1.21244572e+00,  
-9.28660713e-01, 3.63209761e-01, -1.81327713e+00,  
-1.97910760e-01, 4.32874592e-01, -5.40565788e-01,  
2.95630655e-01, 2.07984424e+00, -1.84243292e+00,  
-5.11166017e-01, -5.80935128e-01, 1.06273647e+00,  
1.80551097e-02, 2.27983997e-02, -1.67979148e+00,  
8.12423001e-01, -6.25888706e-01, -1.25027082e+00,  
6.15135458e-01, -1.21394611e-01, -1.28089527e+00,  
3.24609806e-01, 6.70569391e-01, 1.49161323e-01,  
8.01573609e-01, 6.43116741e-01, -9.37629233e-02,  
1.74677366e+00, 6.80996008e-01, -7.03717611e-01,  
1.02079749e-01, 1.19505614e+00, -2.77959386e-01,  
-5.21652916e-02, 3.53683601e-01, -4.08310762e-01,  
-1.81042967e+00, -9.03308062e-01, 1.05404509e+00,  
-9.80876877e-01, 3.52078891e-01, 6.65981840e-01,  
1.06550150e+00, -2.28433613e-01, 3.64483904e-01,  
-1.51484666e+00, -7.52612872e-02, -2.97058082e-01,  
-7.27414382e-01, -2.45875340e-01, -1.27948942e-01,  
-3.69310620e-01, -2.62300428e+00, 2.11585073e+00,  
6.85561585e-01, -1.57443985e-01, 1.38128777e+00,  
6.84265587e-02, 3.12536292e-01, 4.54253185e-01,  
-7.88471875e-01, -6.58403343e-02, -1.41847985e+00,  
-1.39753340e-01, -5.55354856e-01, -5.01917779e-01,  
6.93118522e-01, -2.45360497e-01, -1.26943186e+00,  
-2.62323855e-01))
```

```
n [3]: x_test[0]
```

# Why Overfitting?

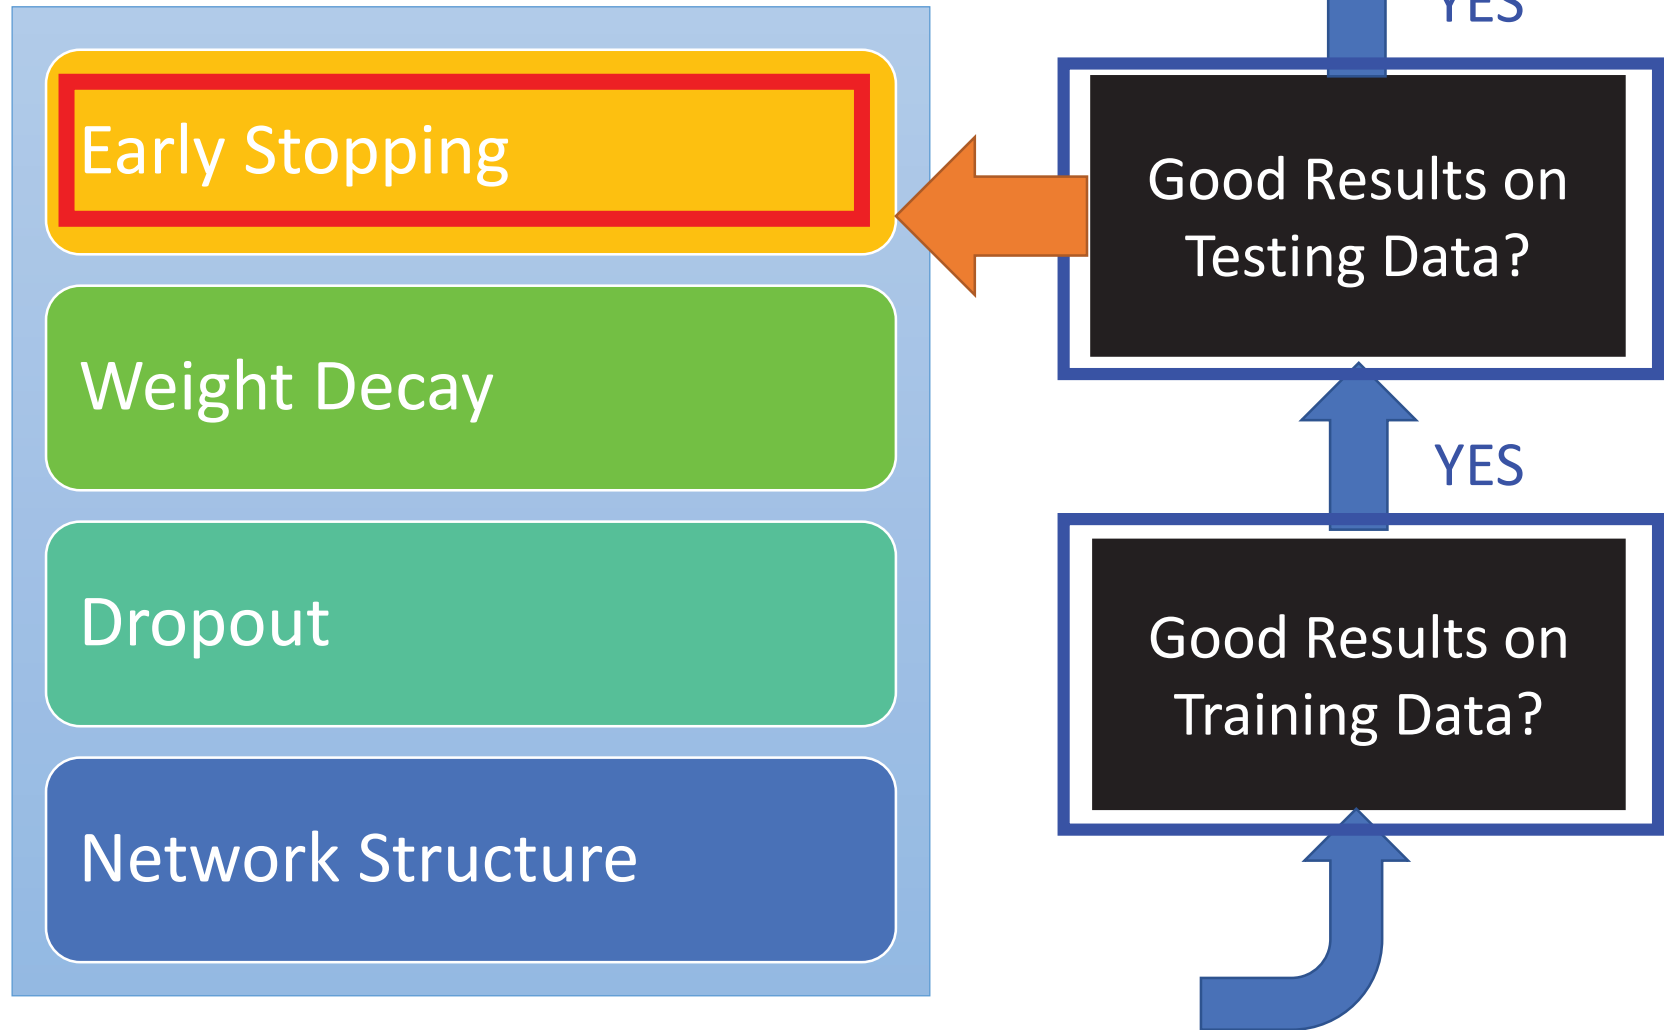
- For experiments, we added some noises to the testing data

Testing:

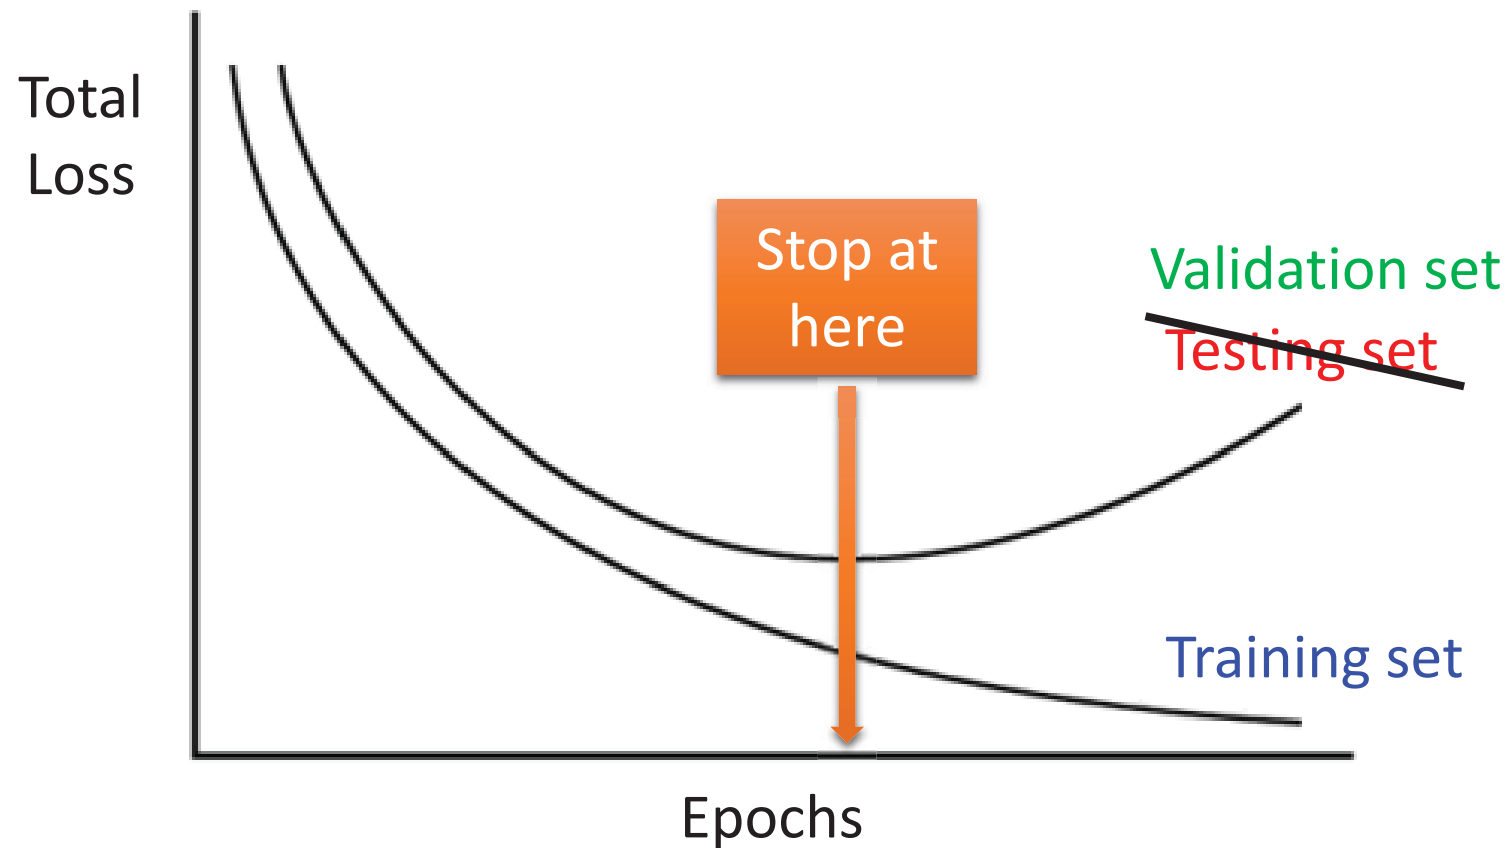
	Accuracy
Clean	0.97
Noisy	0.50

Training is not influenced.

# Recipe of Deep Learning

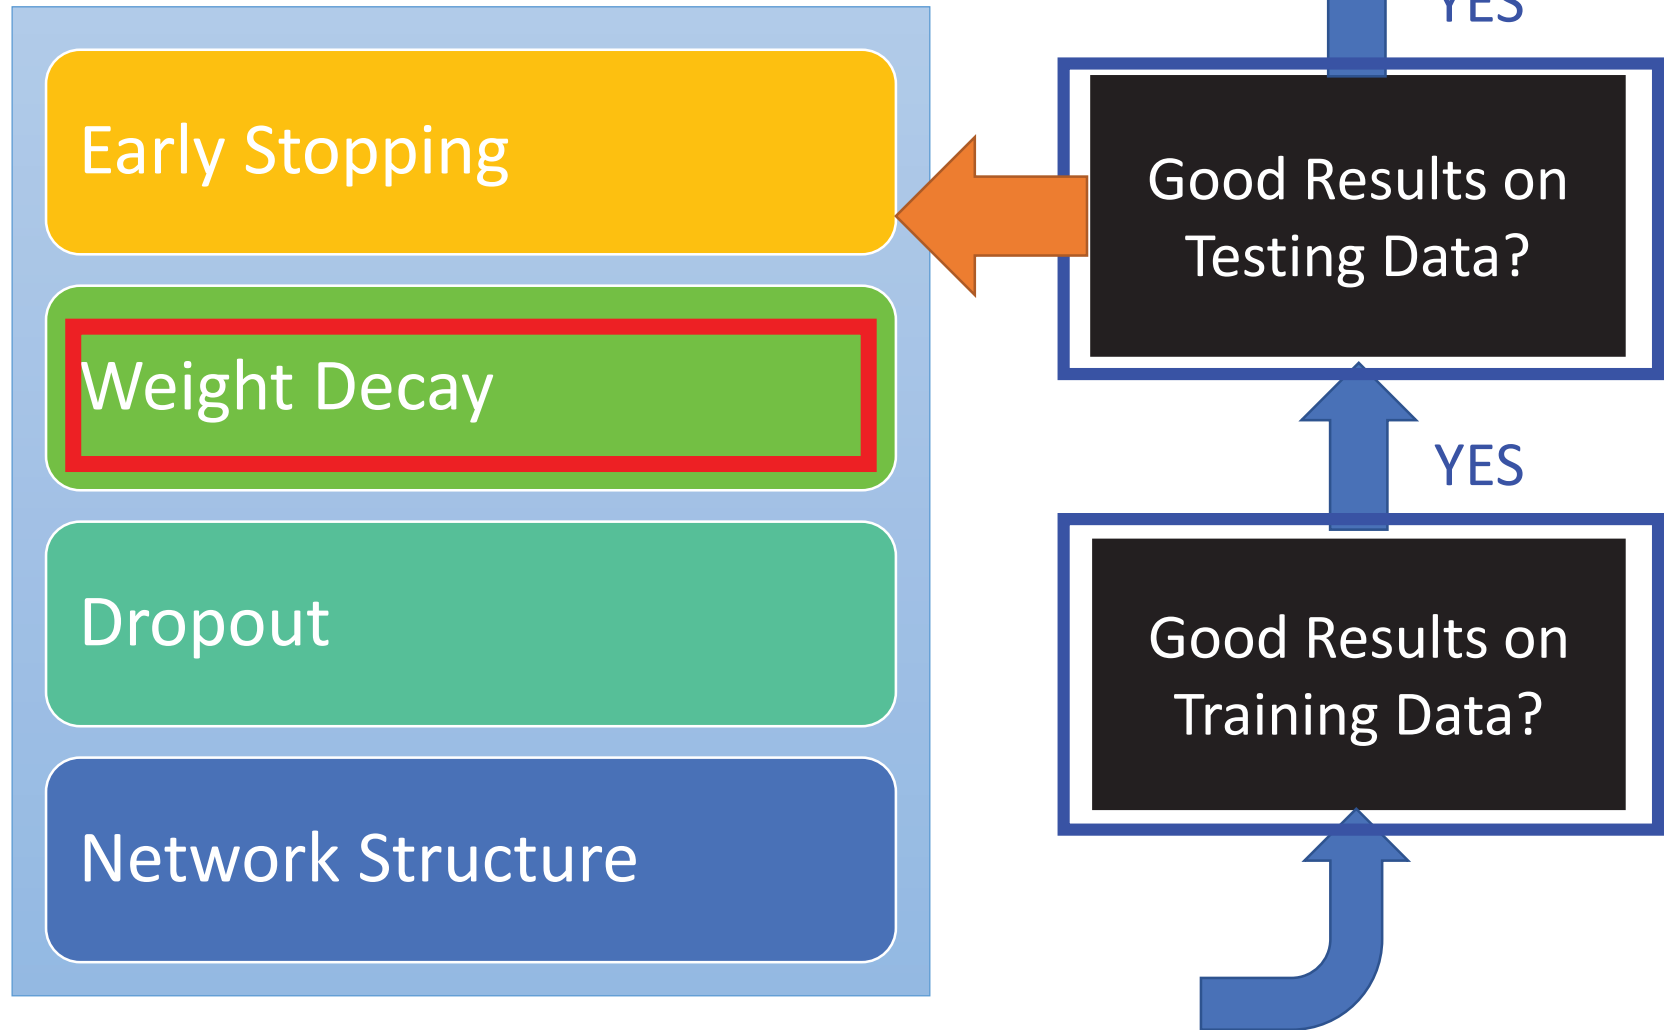


# Early Stopping



**Keras:** <http://keras.io/getting-started/faq/#how-can-i-interrupt-training-when-the-validation-loss-isnt-decreasing-anymore>


# Recipe of Deep Learning





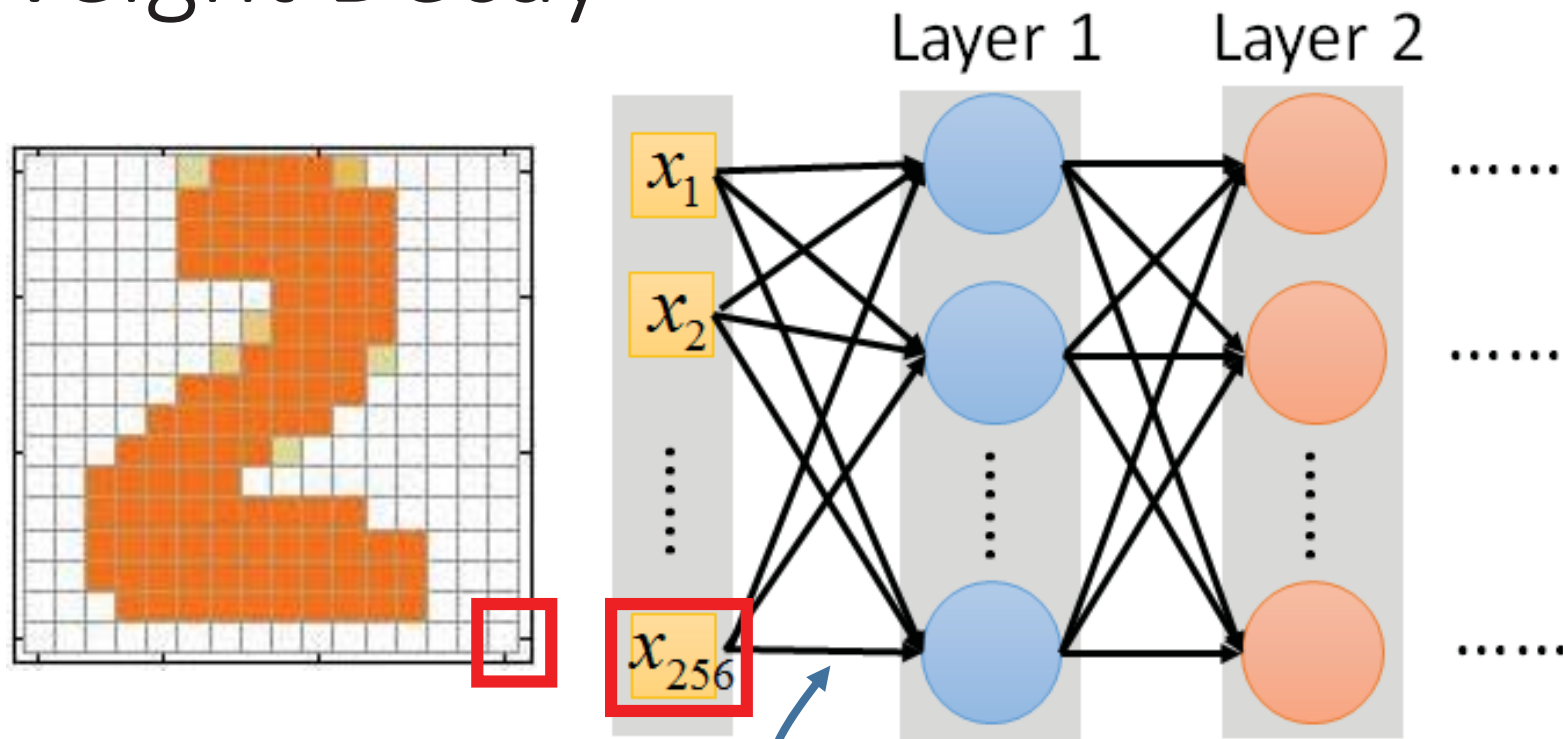
# Weight Decay

- Our brain prunes out the useless link between neurons.



Doing the same thing to machine's brain improves the performance.

# Weight Decay



Weight decay is one kind of regularization

Useless

Close to zero

# Weight Decay

- Implementation

Original:  $w \leftarrow w - \eta \frac{\partial L}{\partial w}$

$$\lambda = 0.01$$

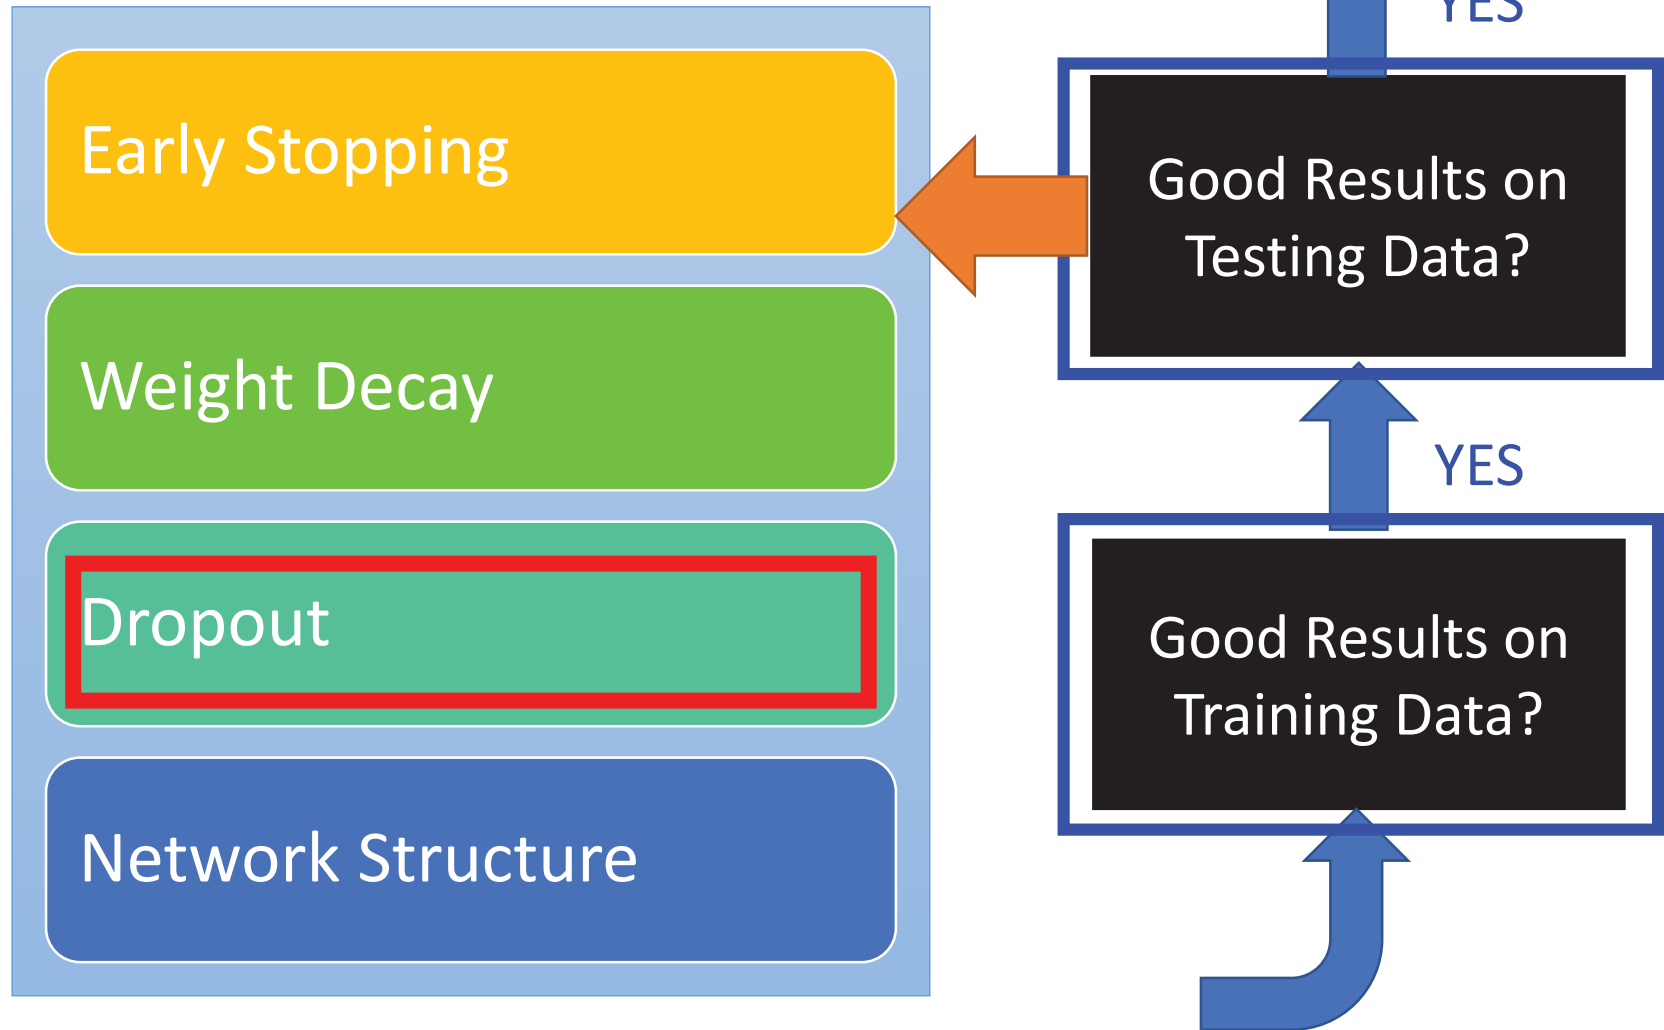
Weight Decay:

$$w \leftarrow \underbrace{0.99}_{\downarrow} w - \eta \frac{\partial L}{\partial w}$$

Smaller and smaller


Keras: <http://keras.io/regularizers/>

# Recipe of Deep Learning



# Dropout


Training:



- **Each time before updating the parameters**
  - Each neuron has  $p\%$  to dropout

# Dropout

Training:




- **Each time before updating the parameters**
  - Each neuron has  $p\%$  to dropout
    - ➡ **The structure of the network is changed.**
  - Using the new network for training

For each mini-batch, we resample the dropout neurons

# Dropout


Testing:



## ➤ No dropout

- If the dropout rate at training is  $p\%$ , all the weights times  $(1-p)\%$
- Assume that the dropout rate is 50%.  
If a weight  $w = 1$  by training, set  $w = 0.5$  for testing.

# Dropout - Intuitive Reason



- When teams up, if everyone expect the partner will do the work, nothing will be done finally.
- However, if you know your partner will dropout, you will do better.
- When testing, no one dropout actually, so obtaining good results eventually.




# Dropout - Intuitive Reason

- Why the weights should multiply (1-p)% (dropout rate) when testing?

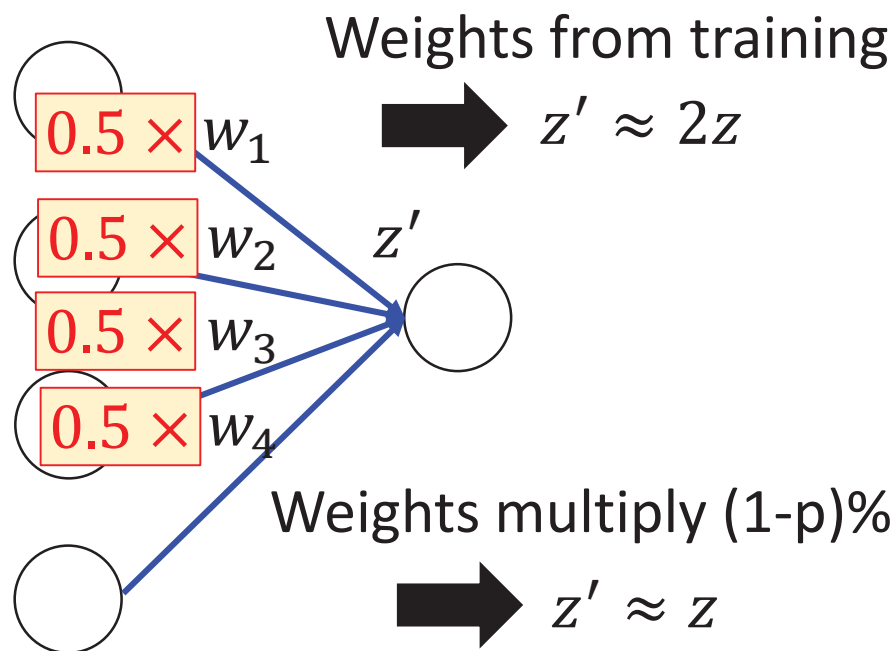
## Training of Dropout

Assume dropout rate is 50%




## Testing of Dropout

No dropout



# Dropout is a kind of ensemble.


**Ensemble**




Train a bunch of networks with different structures

# Dropout is a kind of ensemble.

## Ensemble




# Dropout is a kind of ensemble.



- Using one mini-batch to train one network
- Some parameters in the network are shared

# Dropout is a kind of ensemble.

## Testing of Dropout

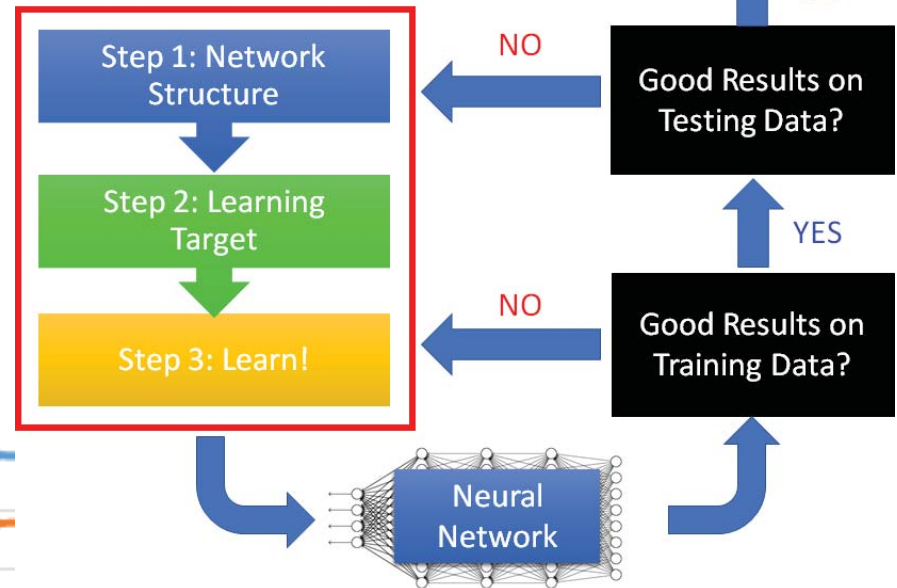



# More about dropout

- More reference for dropout [[Nitish Srivastava, JMLR'14](#)] [[Pierre Baldi, NIPS'13](#)][[Geoffrey E. Hinton, arXiv'12](#)]
- Dropout works better with Maxout [[Ian J. Goodfellow, ICML'13](#)]
- Dropconnect [[Li Wan, ICML'13](#)]
  - Dropout delete neurons
  - Dropconnect deletes the connection between neurons
- Annealed dropout [[S.J. Rennie, SLT'14](#)]
  - Dropout rate decreases by epochs
- Standout [[J. Ba, NIPS'13](#)]
  - Each neural has different dropout rate

# Let's try it

No Dropout



Testing:

	Accuracy
Noisy	0.50
+ dropout	0.63

# Variants of Neural Networks

Convolutional Neural  
Network (CNN)


Widely used in  
image processing

Recurrent Neural Network  
(RNN)



# Why CNN for Image?

- When processing image, the first layer of fully connected network would be very large



Can the fully connected network be simplified by considering the properties of image recognition?

# Why CNN for Image

- Some patterns are much smaller than the whole image

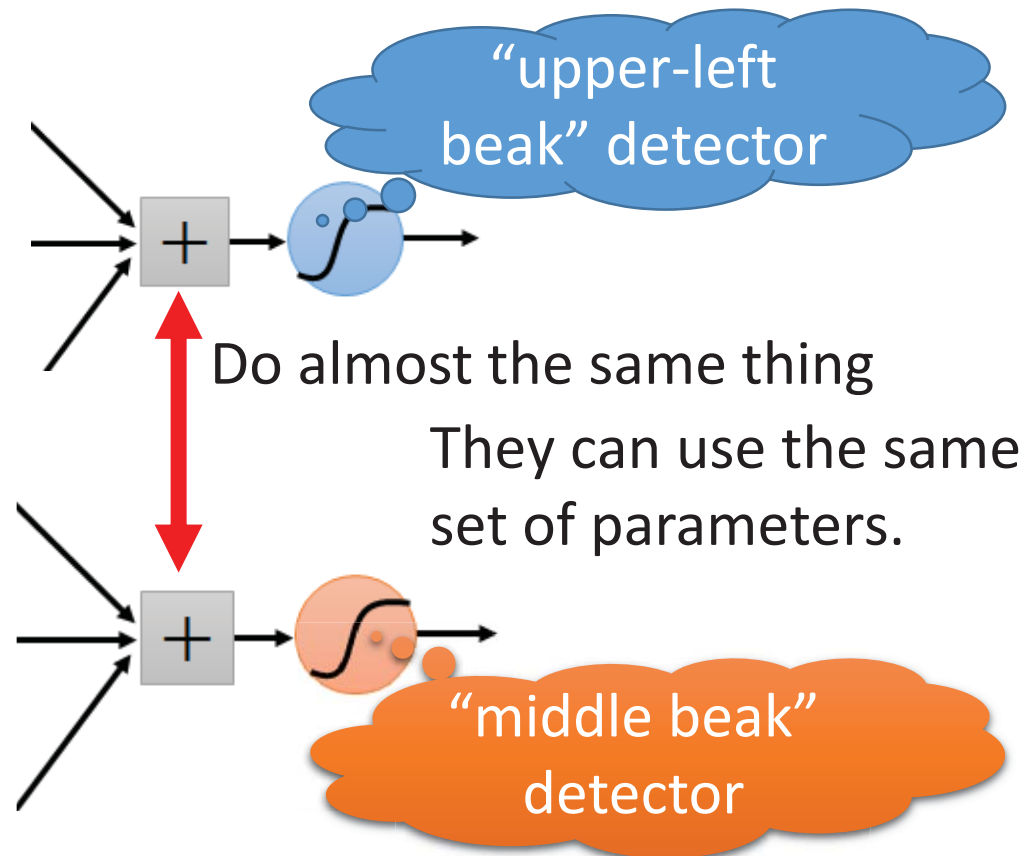

A neuron does not have to see the whole image to discover the pattern.

Connecting to small region with less parameters



# Why CNN for Image

- The same patterns appear in different regions.



# Why CNN for Image

- Subsampling the pixels will not change the object

bird



subsampling

bird

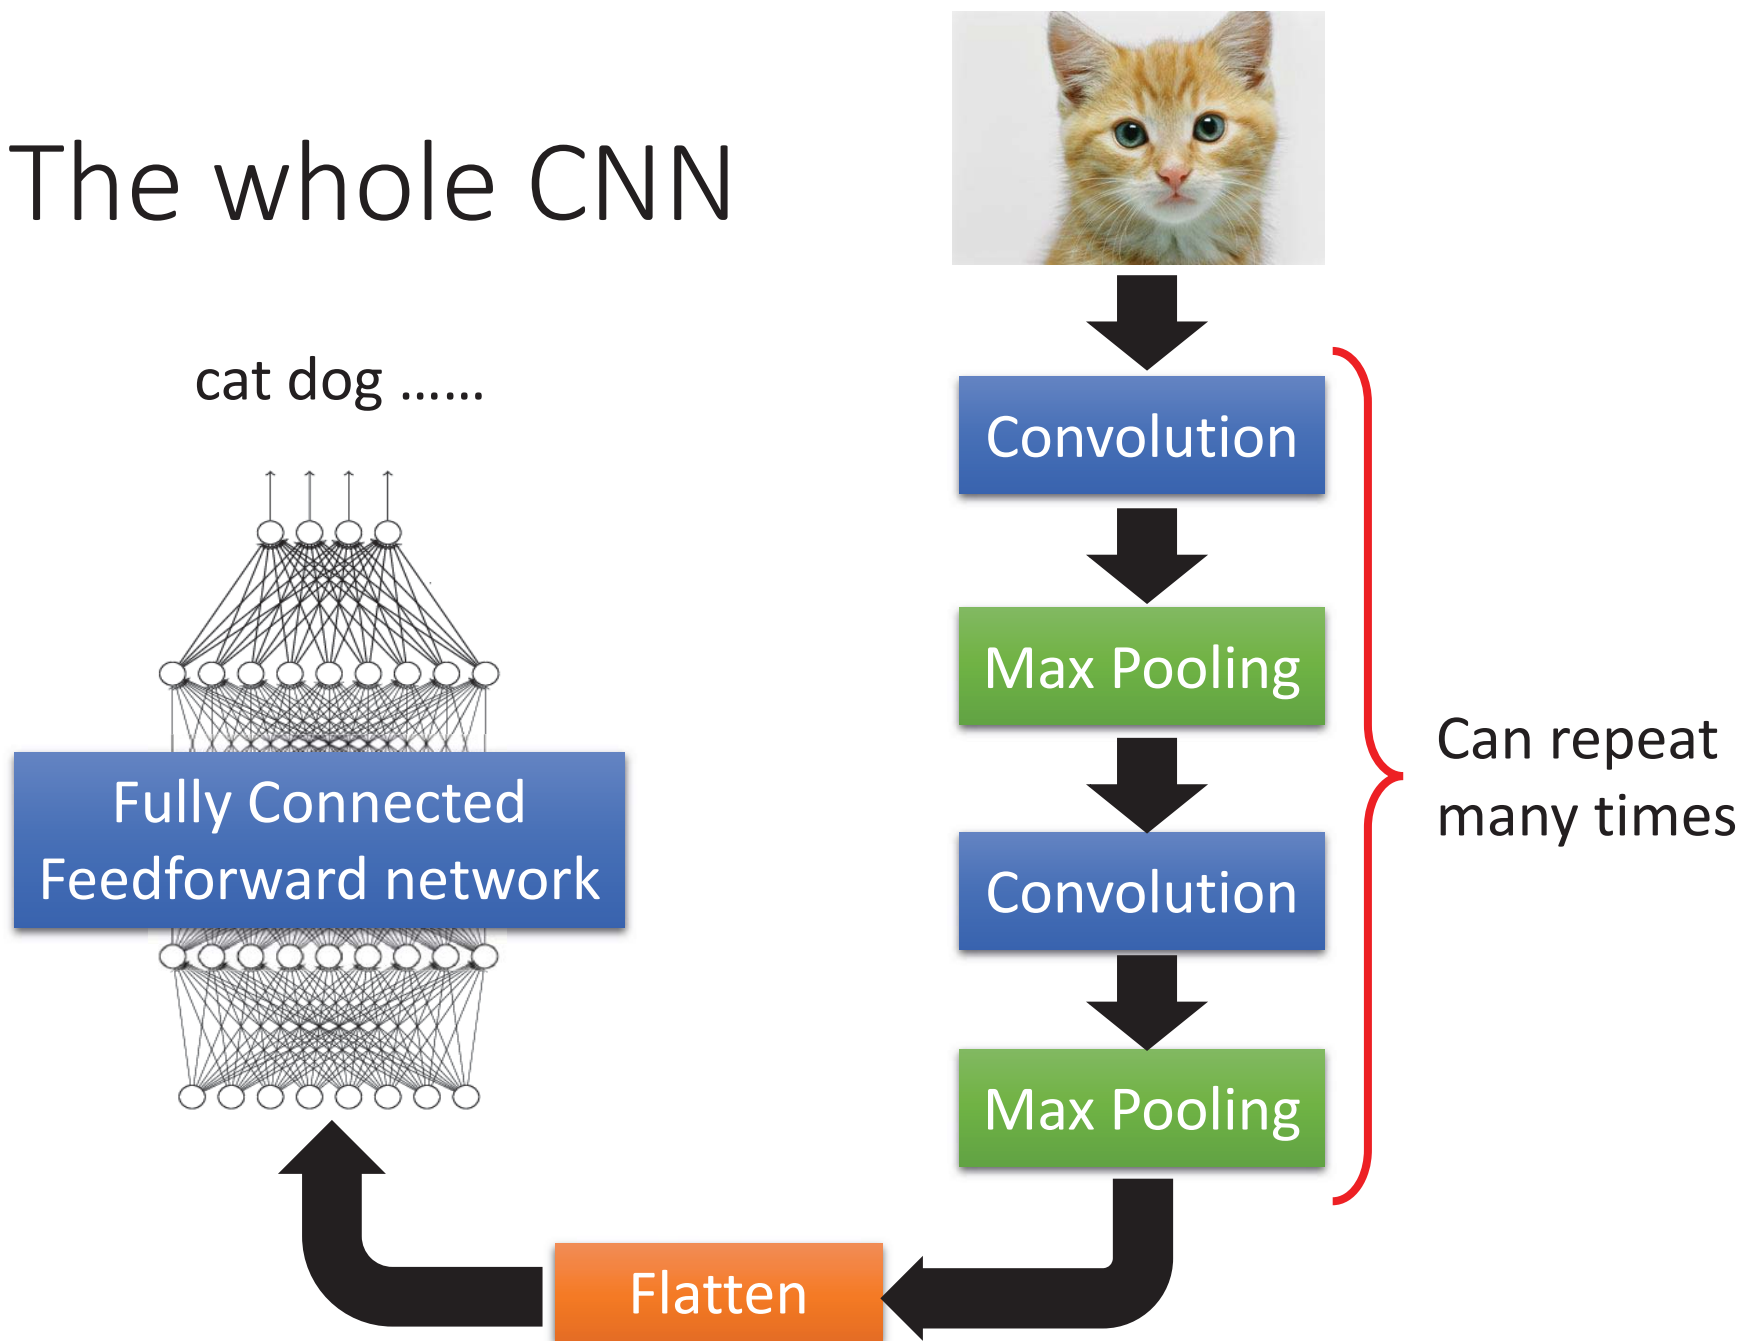


We can subsample the pixels to make image smaller



Less parameters for the network to process the image

# The whole CNN



# The whole CNN

## Property 1

- Some patterns are much smaller than the whole image

## Property 2

- The same patterns appear in different regions.

## Property 3

- Subsampling the pixels will not change the object



Convolution

Max Pooling

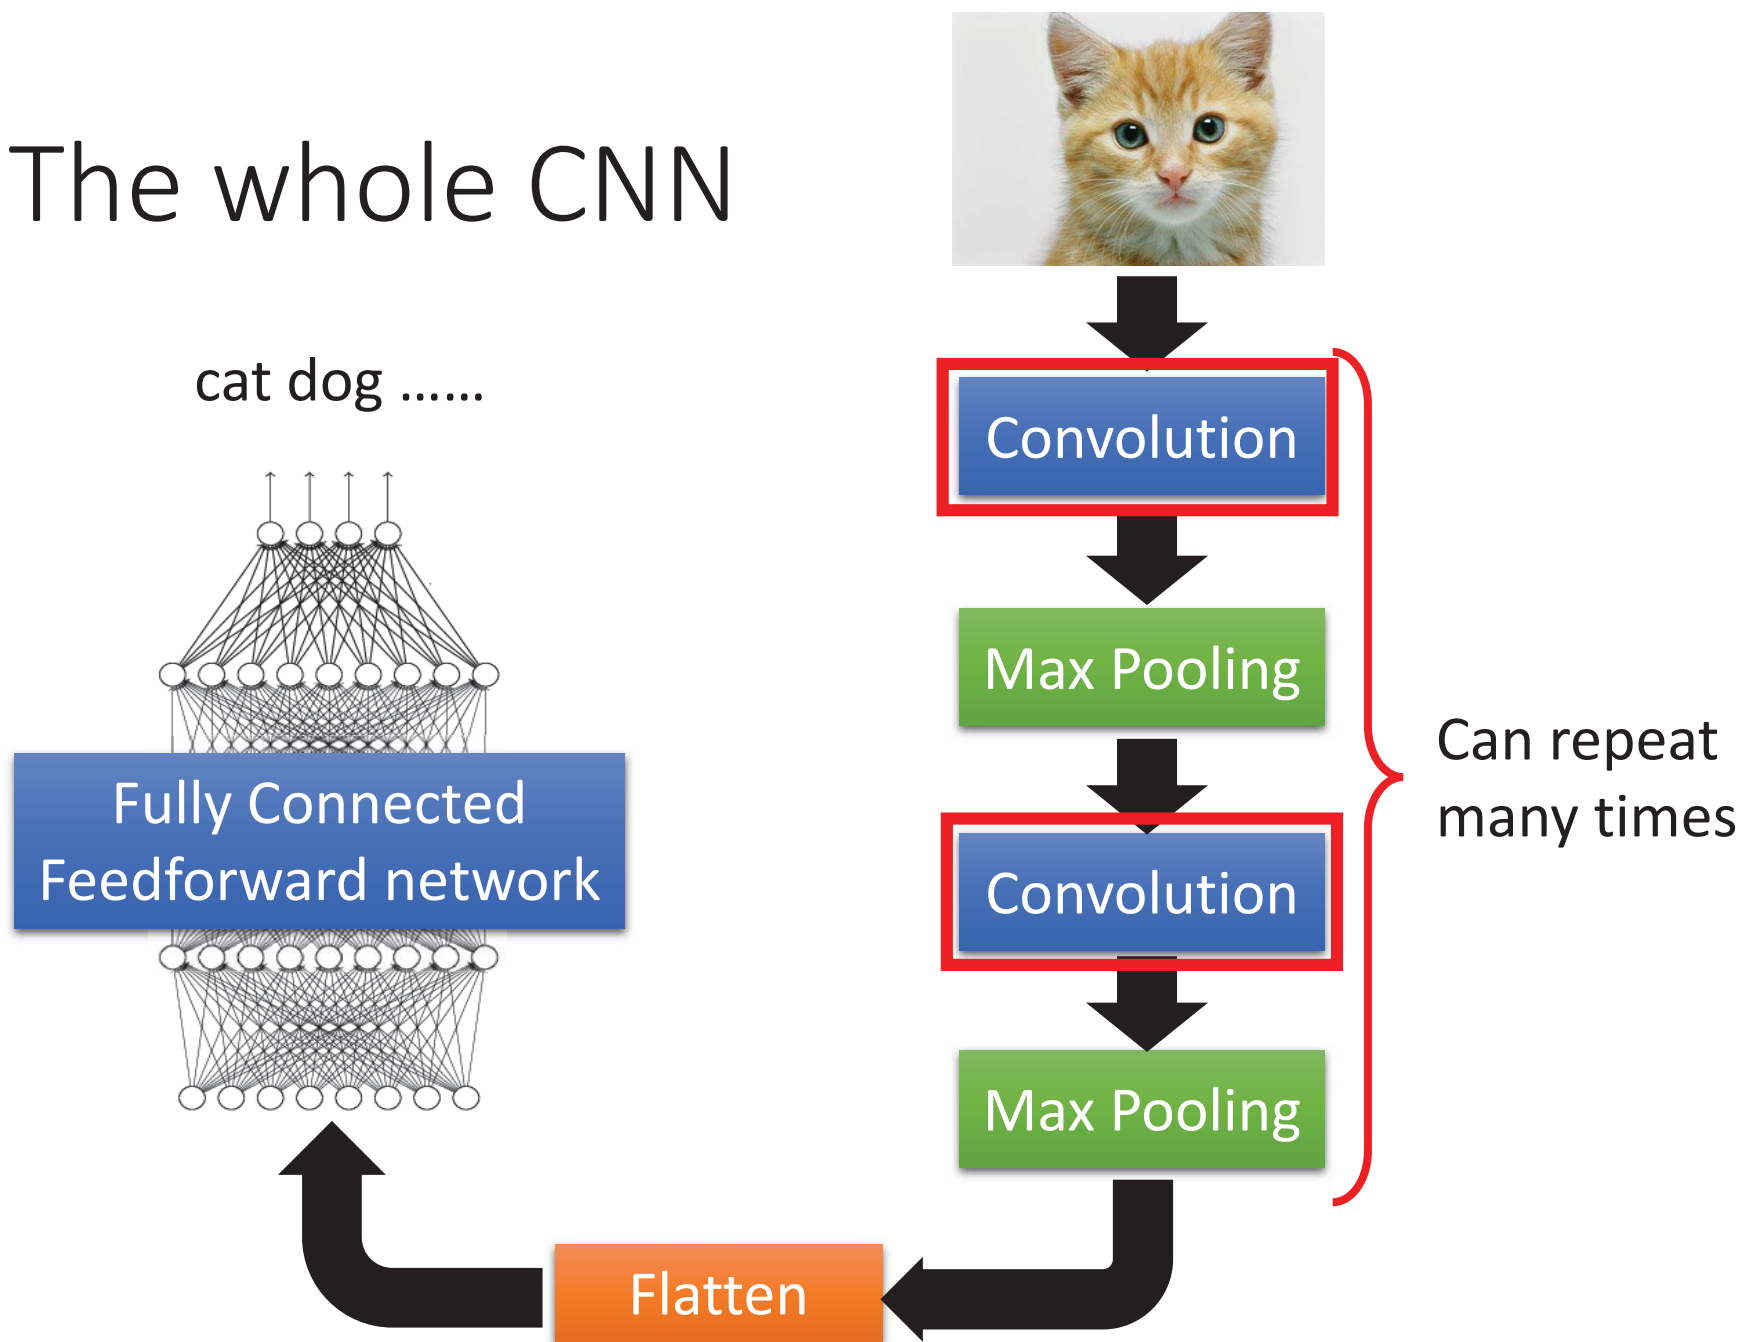
Convolution

Max Pooling

Can repeat many times

Flatten

# The whole CNN



# CNN – Convolution

Those are the network parameters to be learned.

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1  
Matrix

-1	1	-1
-1	1	-1
-1	1	-1

Filter 2  
Matrix

⋮

Property 1

Each filter detects a small pattern (3 x 3).



# CNN – Convolution


1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

stride=1

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image



# CNN – Convolution


1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

If stride=2

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image



We set stride=1 below

# CNN – Convolution

stride=1

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

3	-1	-3	-1
-3	1	0	-3
-3	-3	0	1
3	-2	-2	-1

Property 2

# CNN – Convolution

stride=1


1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

-1	1	-1
-1	1	-1
-1	1	-1

Filter 2

Do the same process for every filter



4 x 4 image

# CNN – Zero Padding

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

0	0	0				
0	1	0	0	0	0	1
0	0	1	0	0	1	0
	0	0	1	1	0	0
	1	0	0	0	1	0
	0	1	0	0	1	0
	0	0	1	0	1	0

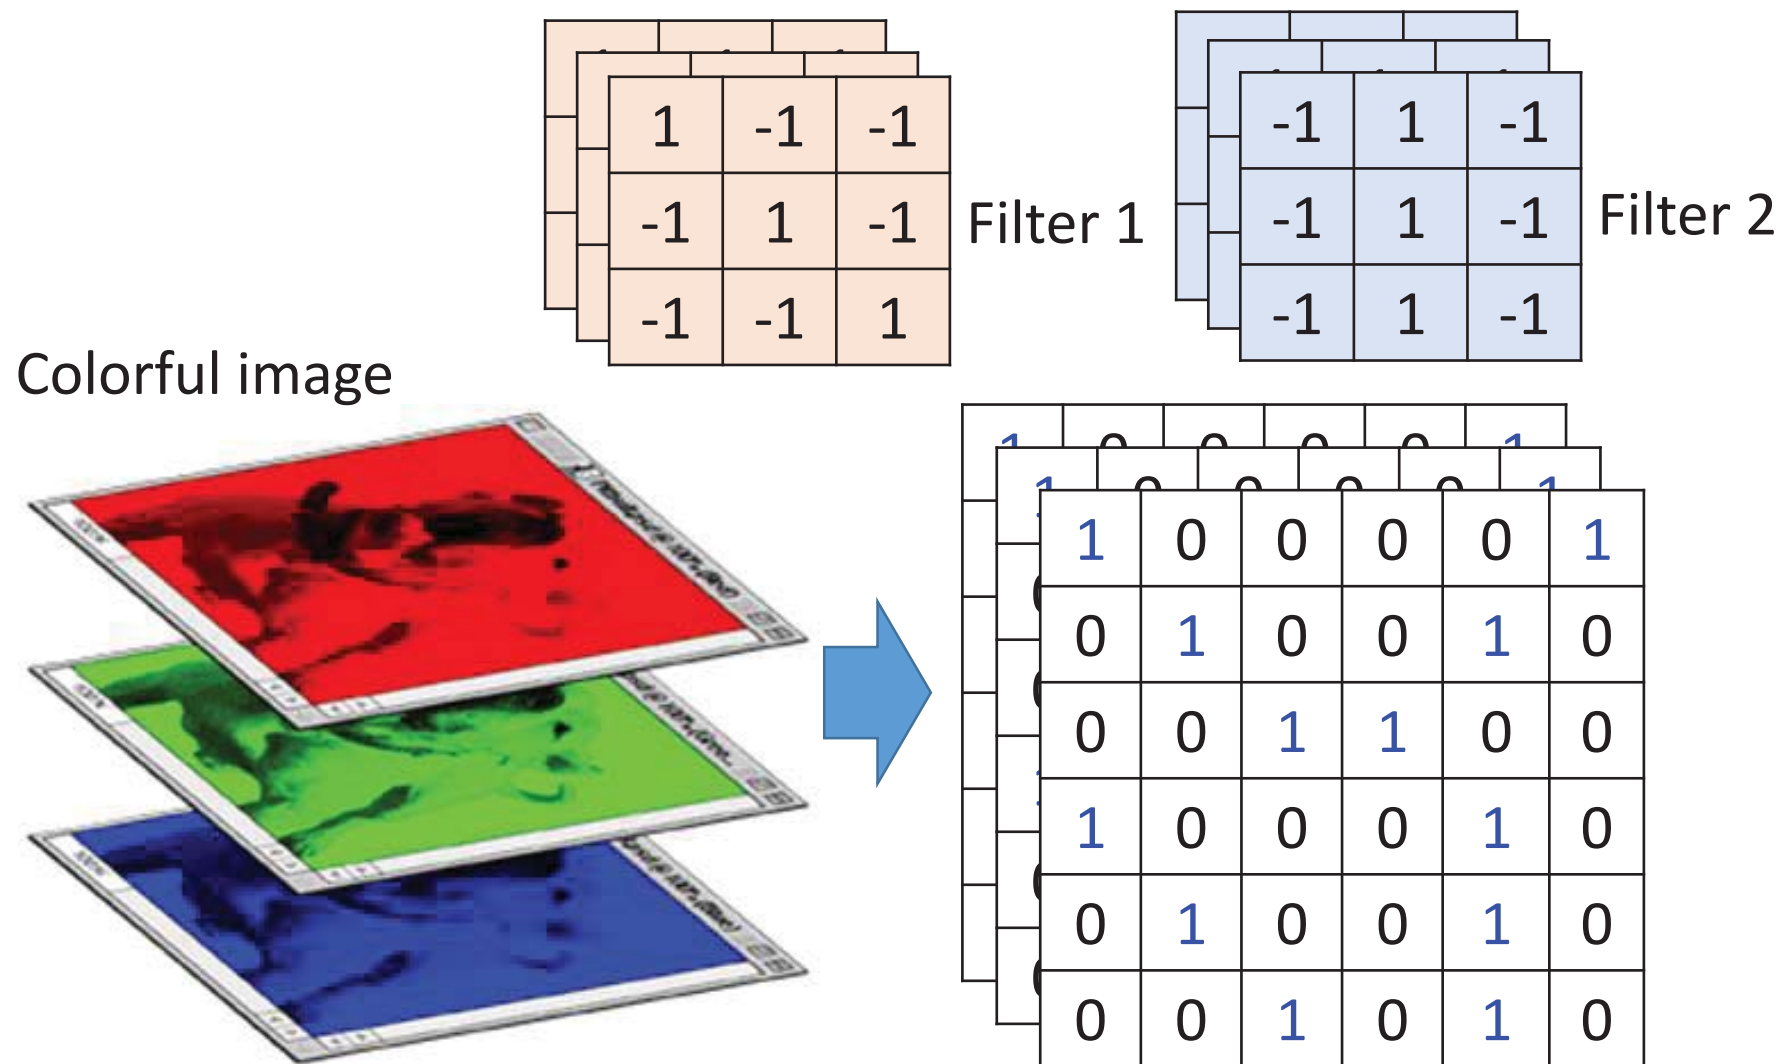
6 x 6 image

You will get another 6 x 6 images in this way

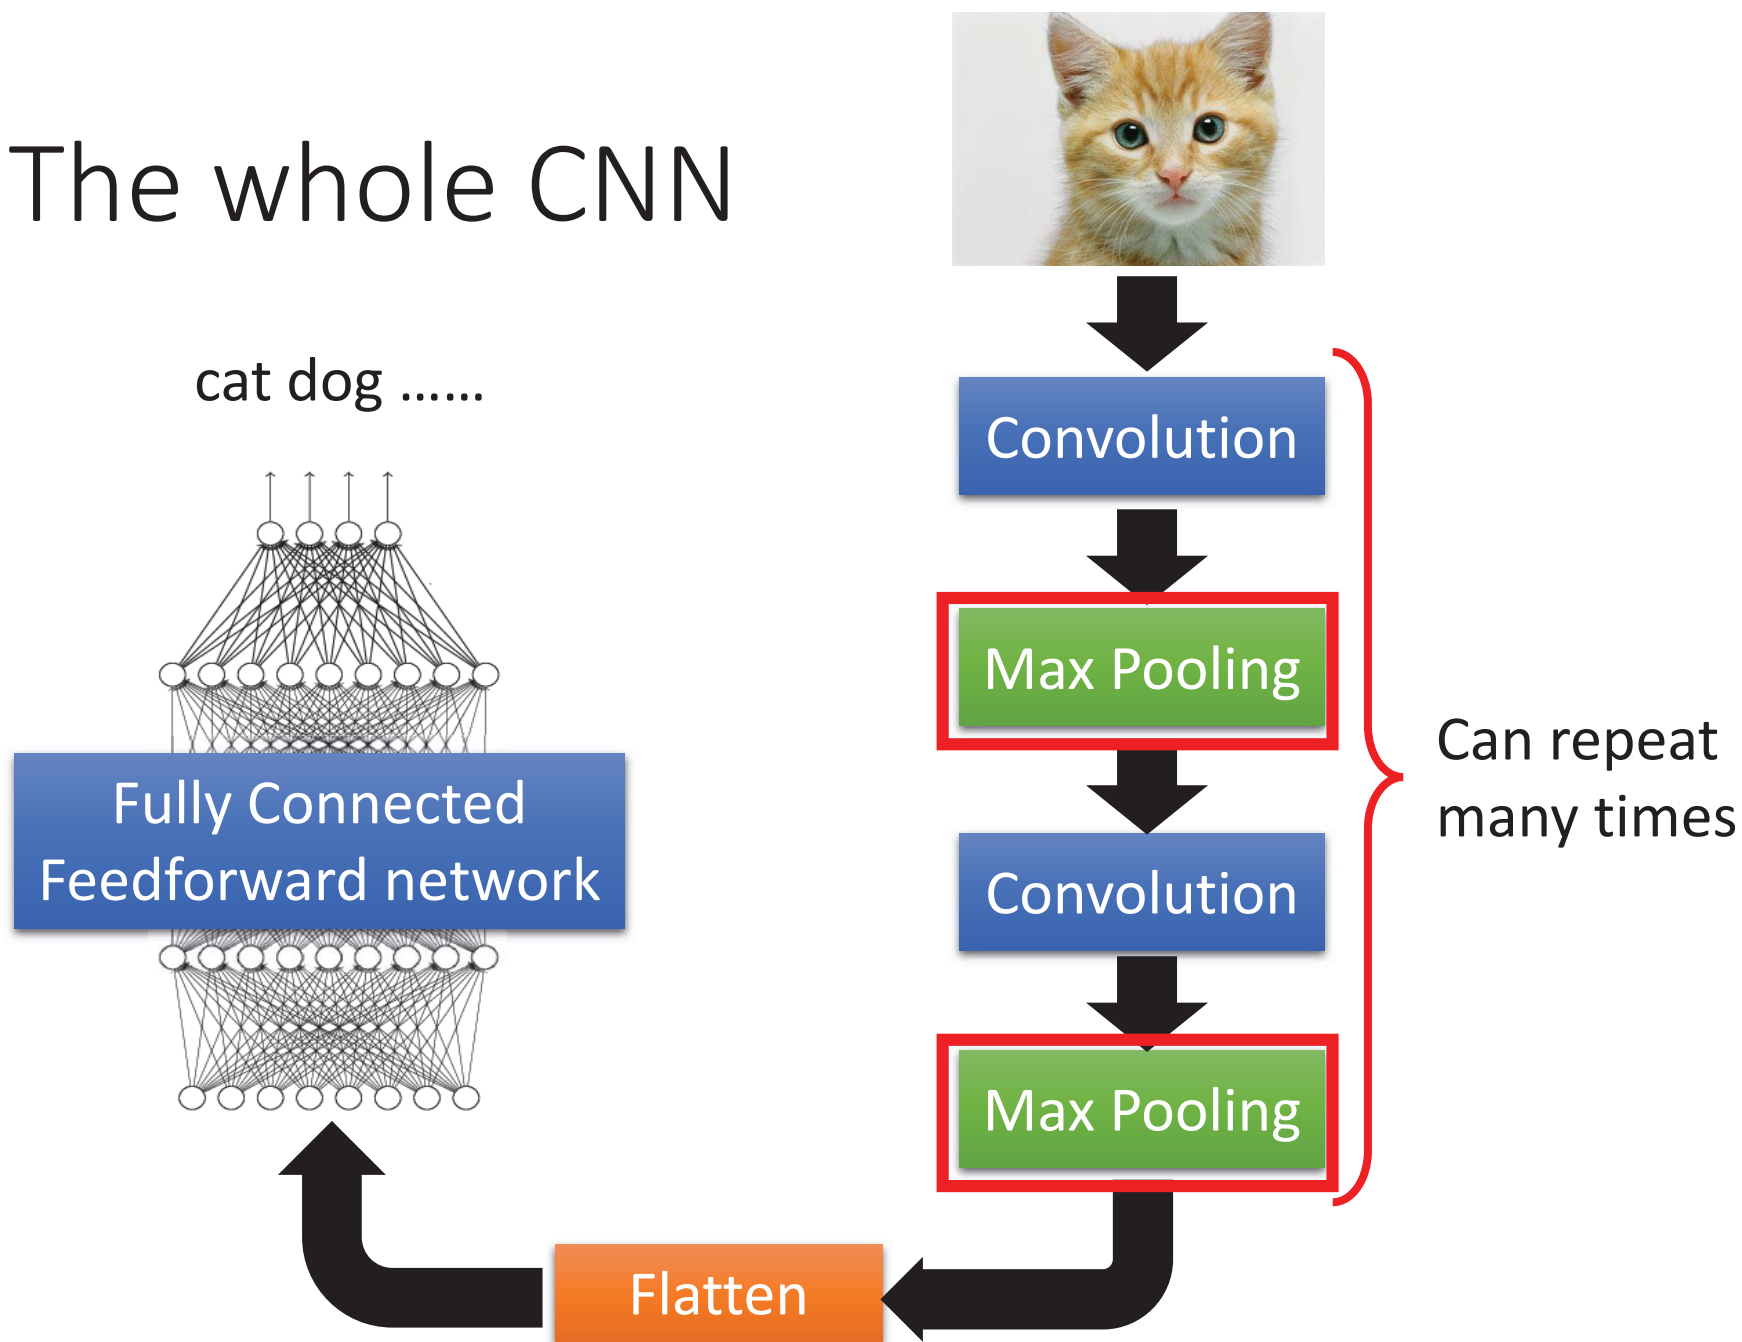


Zero padding

# CNN – Colorful image



# The whole CNN



# CNN – Max Pooling

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

-1	1	-1
-1	1	-1
-1	1	-1

Filter 2

3	-1	-3	-1
-3	1	0	-3
-3	-3	0	1
3	-2	-2	-1


-1	-1	-1	-1
-1	-1	-2	1
-1	-1	-2	1
-1	0	-4	3



# CNN – Max Pooling

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image




Conv



Max  
Pooling




New image  
but smaller



2 x 2 image

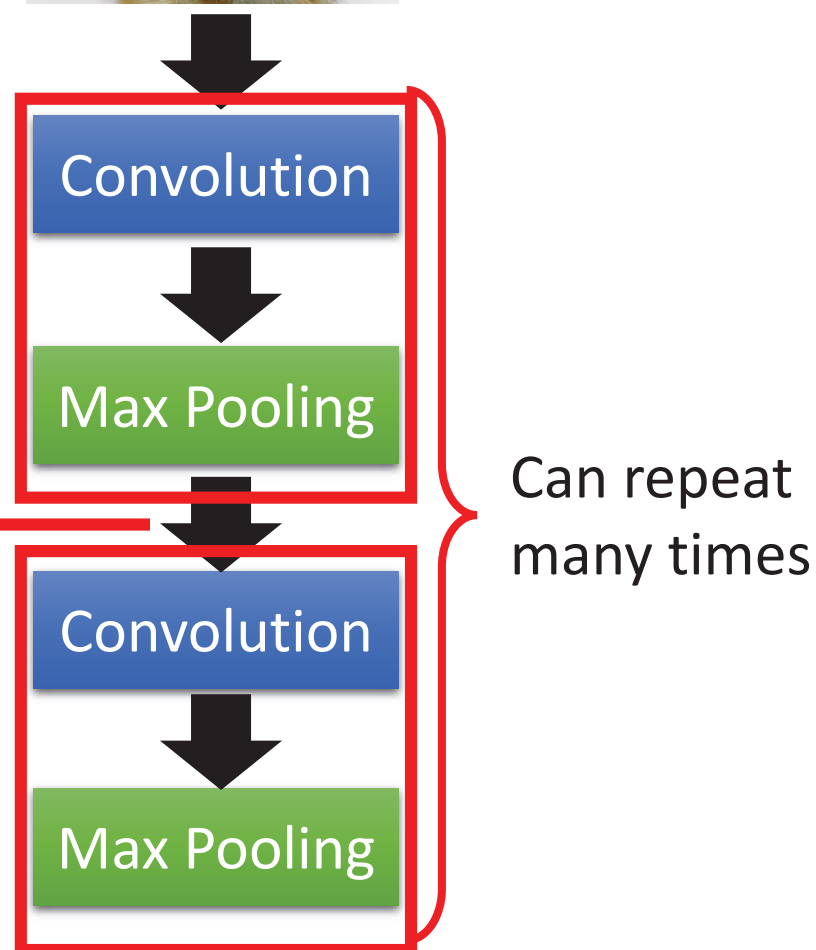
Each filter  
is a channel

# The whole CNN




Smaller than the original image

The number of the channel is the number of filters



# The whole CNN

cat dog .....



Convolution

Max Pooling

A new image

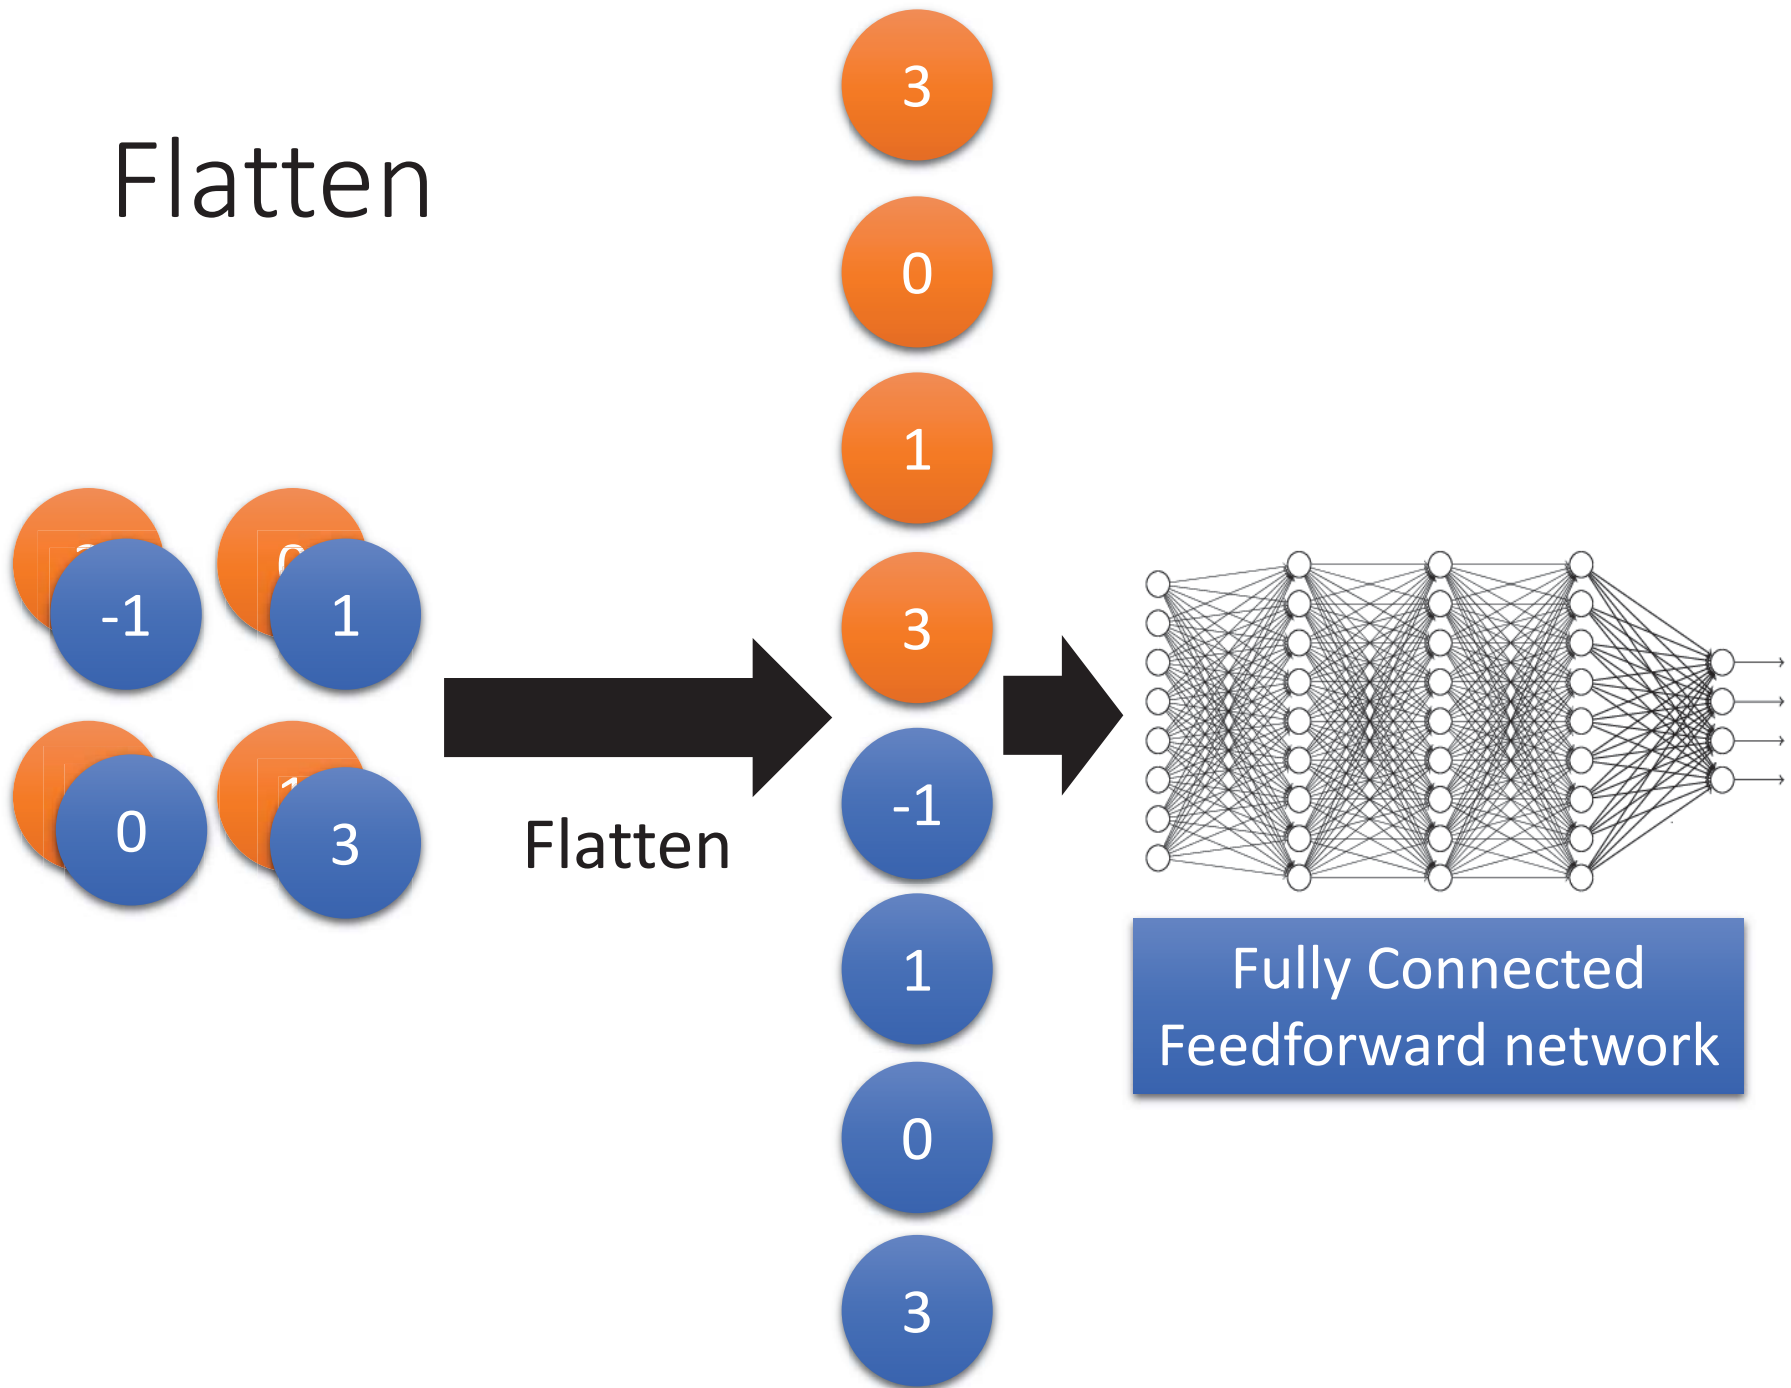
Convolution

Max Pooling



A new image

Flatten

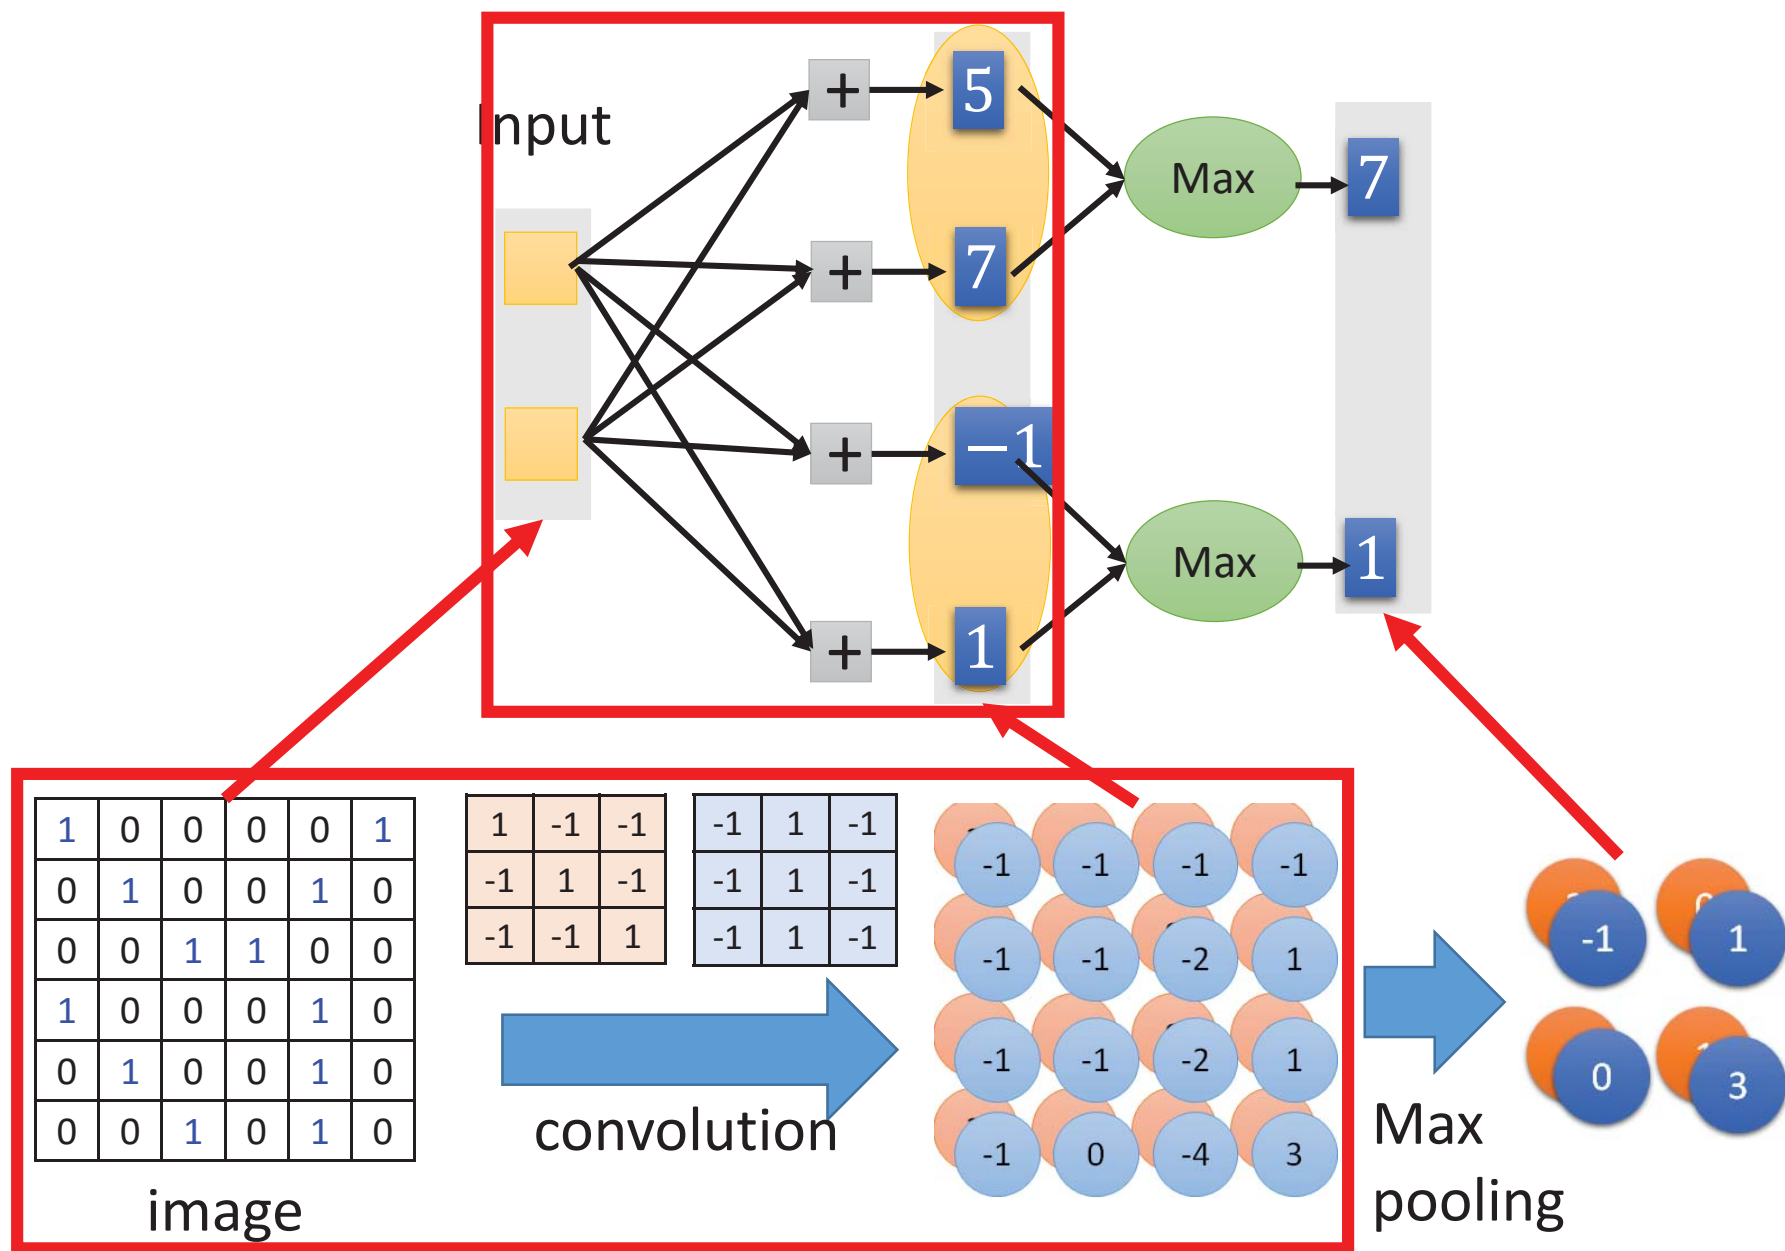
# Flatten




# The whole CNN




Can repeat many times




(Ignoring the non-linear activation function after the convolution.)




Less parameters!



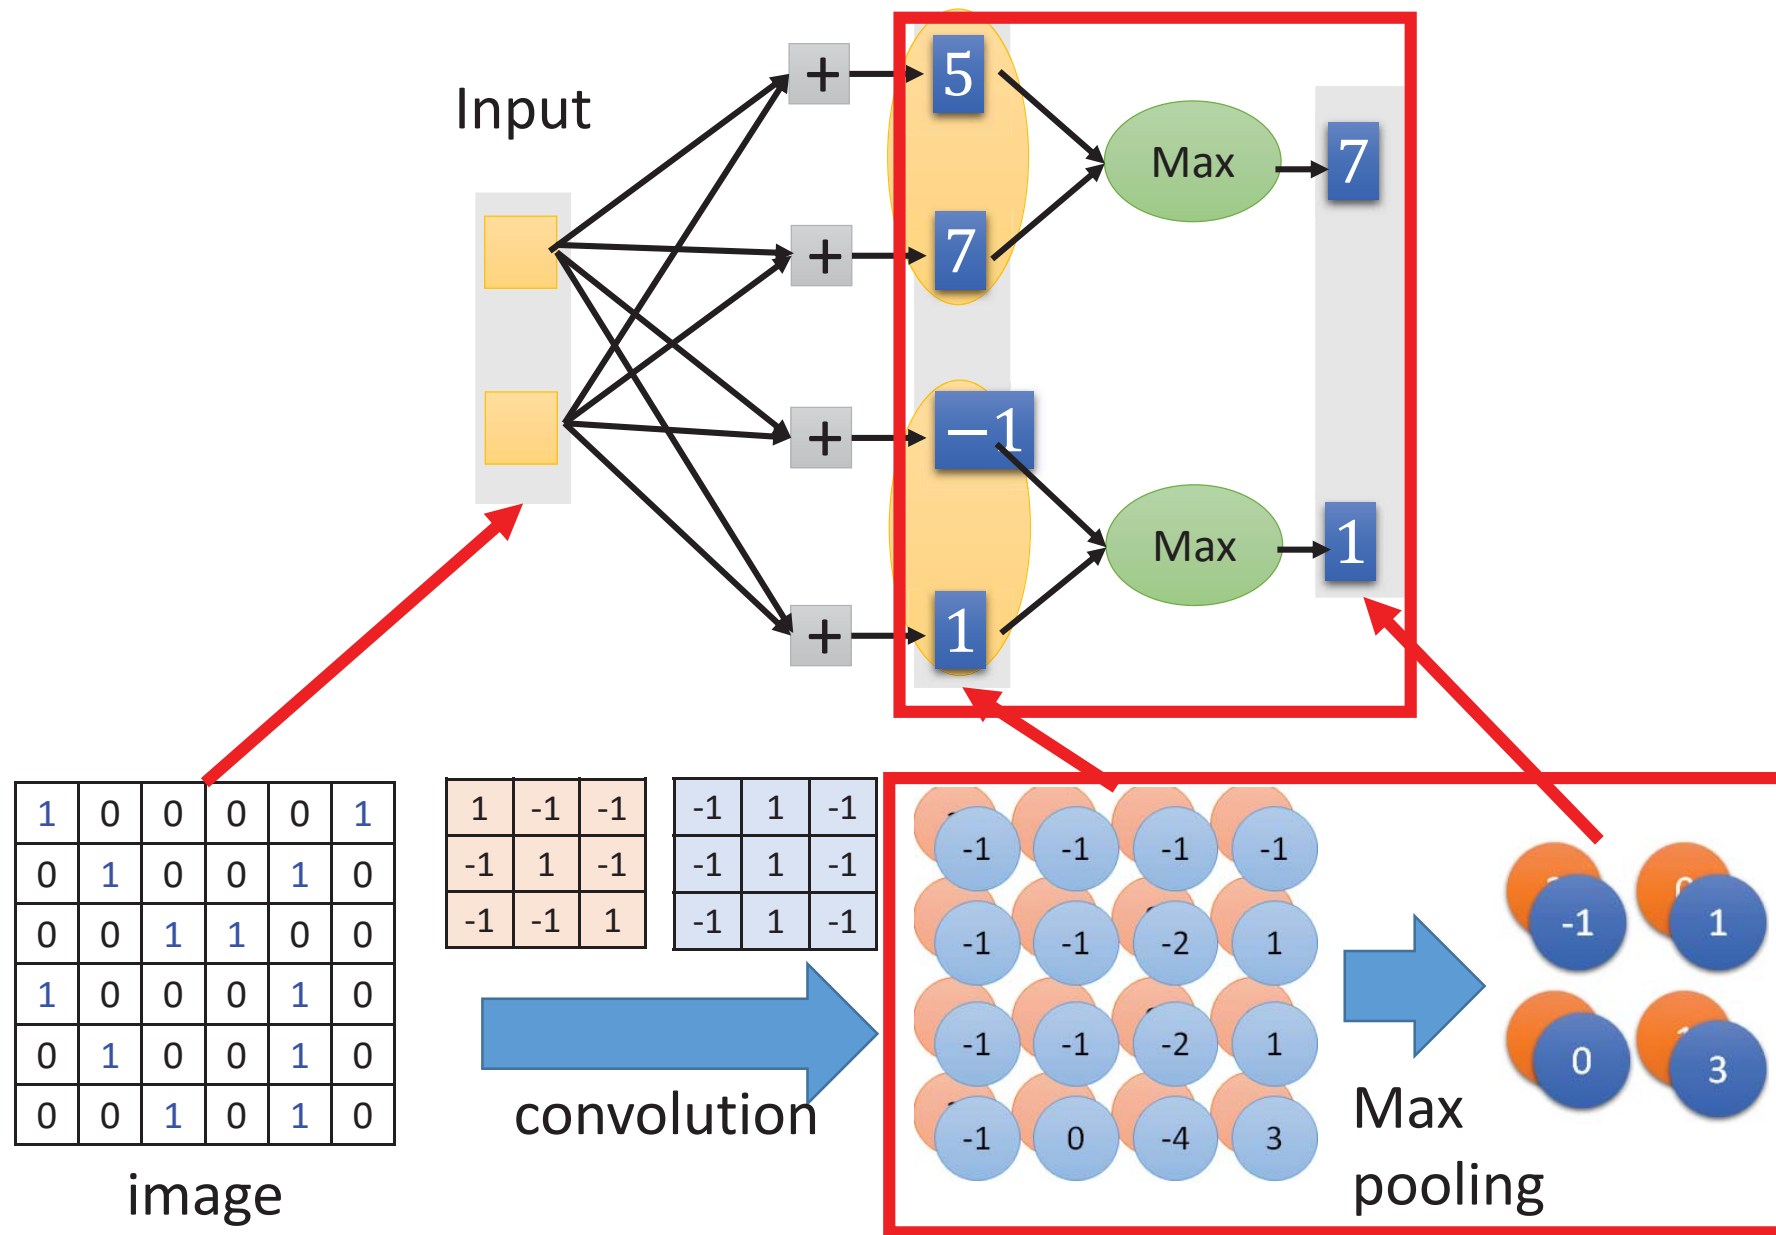


Less parameters!



Even less parameters!

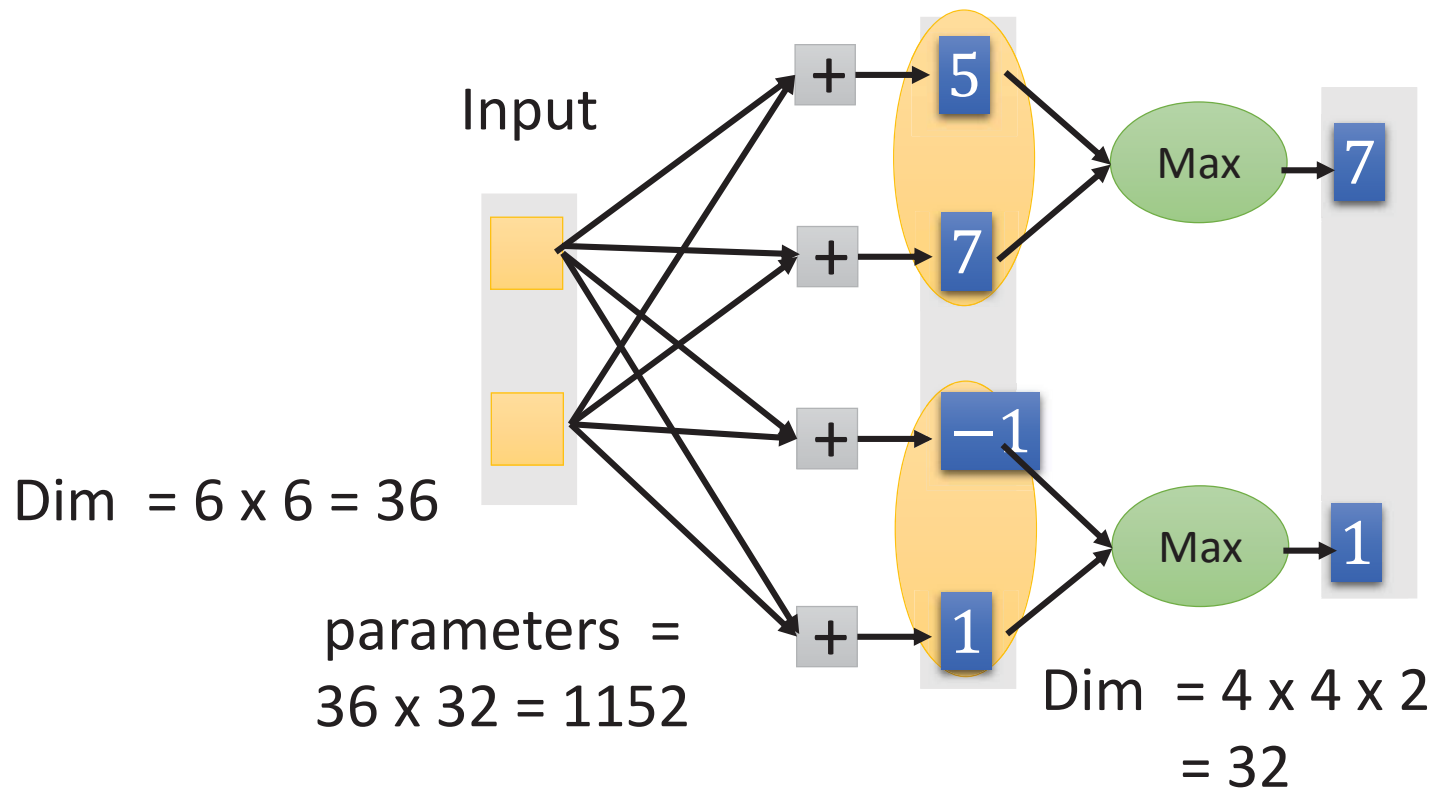






(Ignoring the non-linear activation function after the convolution.)






1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0



image

1	-1	-1
-1	1	-1
-1	-1	1


-1	1	-1
-1	1	-1
-1	1	-1



Only  $9 \times 2 = 18$   
parameters



# Convolutional Neural Network



Learning: Nothing special, just gradient descent .....

# Playing Go



19 x 19 matrix  
(image)

Black: 1  
white: -1  
none: 0



Network



Next move  
(19 x 19  
positions)

19 x 19 vector

Fully-connected feedword  
network can be used

But CNN performs much better.

# Variants of Neural Networks

Convolutional Neural  
Network (CNN)

Recurrent Neural Network  
(RNN)

Neural Network with Memory

# Example Application

- Slot Filling

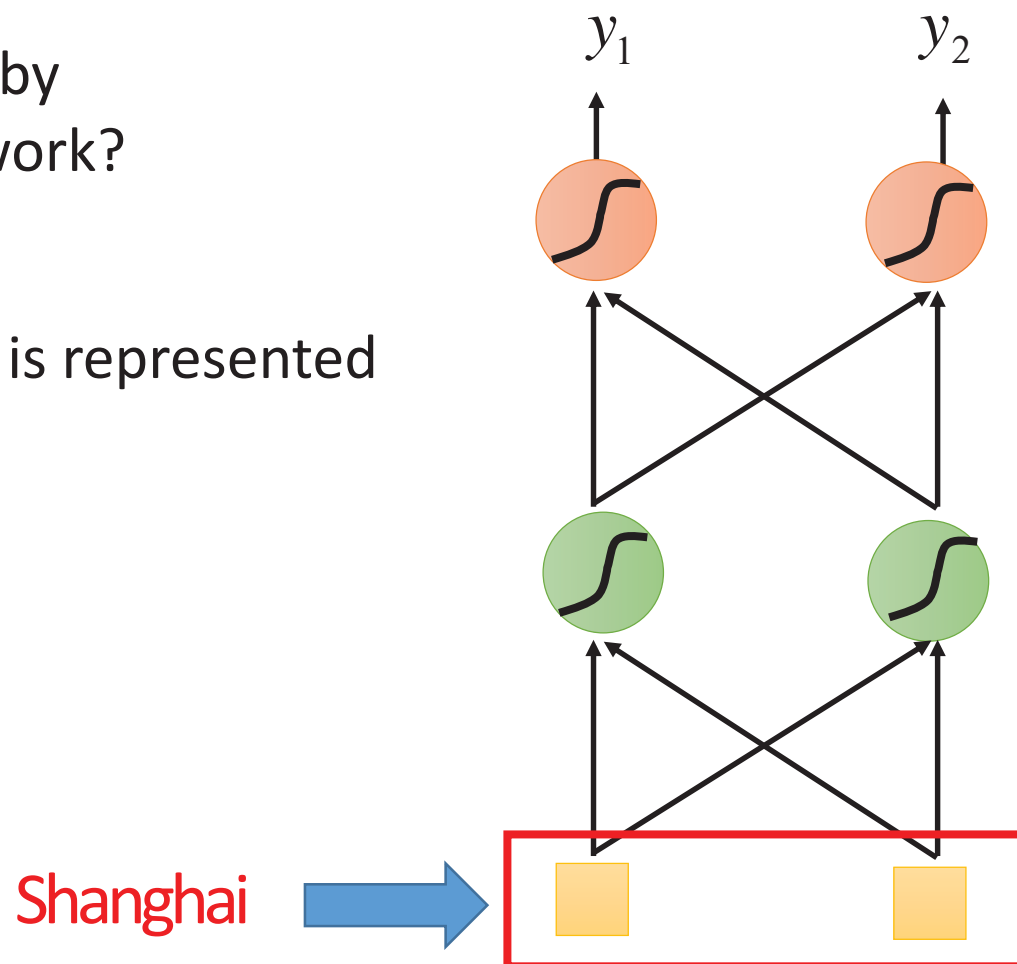


# Example Application

Solving slot filling by  
Feedforward network?

Input: a word

(Each word is represented  
as a vector)





# 1-of-N encoding

How to represent each word as a vector?

**1-of-N Encoding**    lexicon = {apple, bag, cat, dog, elephant}

The vector is lexicon size.

Each dimension corresponds  
to a word in the lexicon

The dimension for the word  
is 1, and others are 0

apple = [ 1   0   0   0   0 ]

bag    = [ 0   1   0   0   0 ]


cat    = [ 0   0   1   0   0 ]

dog    = [ 0   0   0   1   0 ]


elephant = [ 0   0   0   0   1 ]

# Beyond 1-of-N encoding

## Dimension for "Other"



## Word hashing



# Example Application

Solving slot filling by  
Feedforward network?


Input: a word

(Each word is represented  
as a vector)


Output:

Probability distribution that  
the input word belonging to  
the slots

Shanghai



# Example Application




Problem?




Neural network  
needs memory!

Shanghai →



# Recurrent Neural Network (RNN)

The output of hidden layer are stored in the memory.



Memory can be considered as another input.

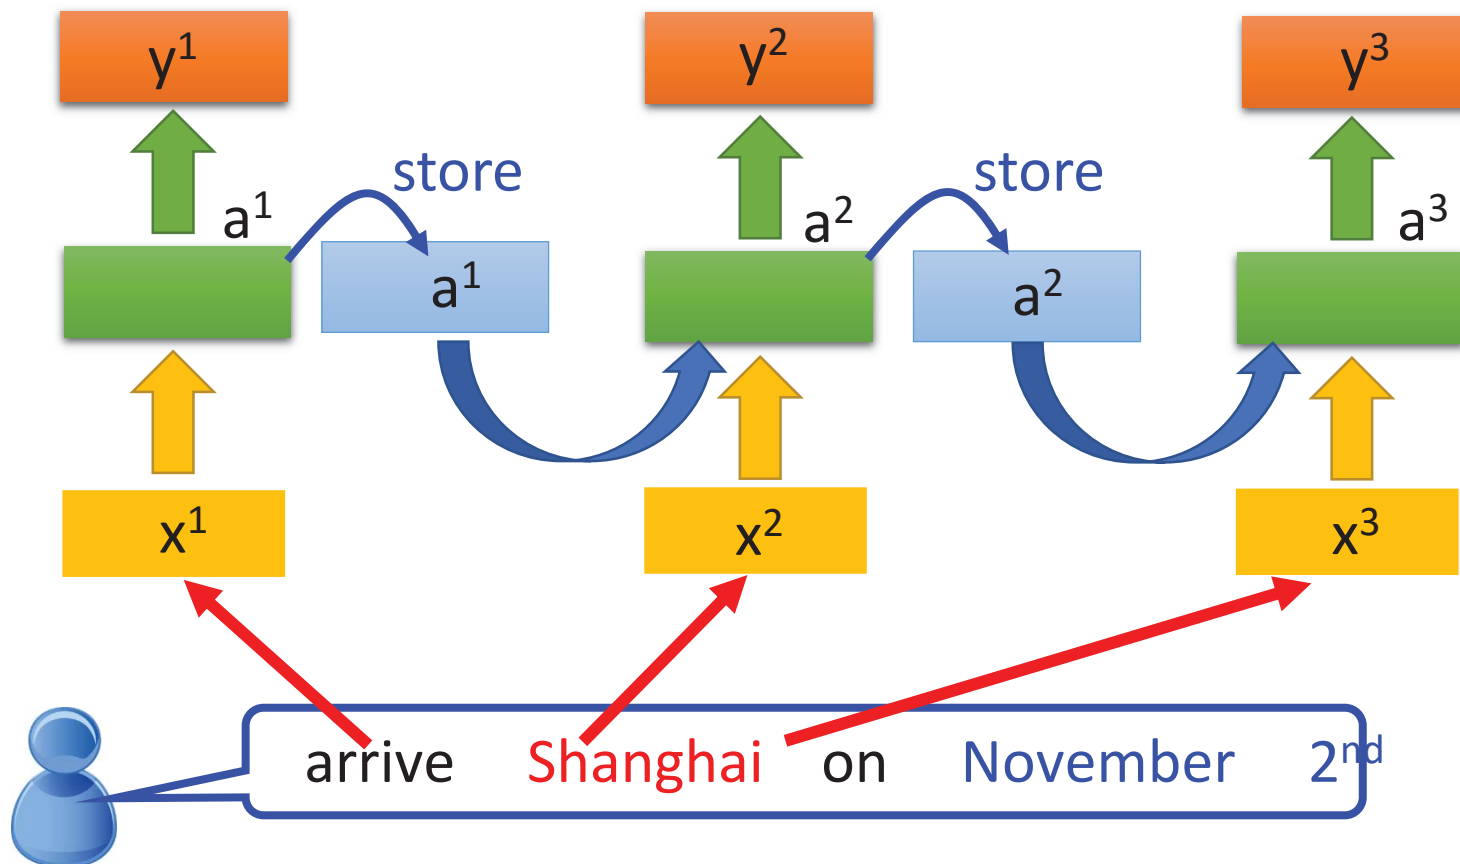
# RNN

The same network is used again and again.

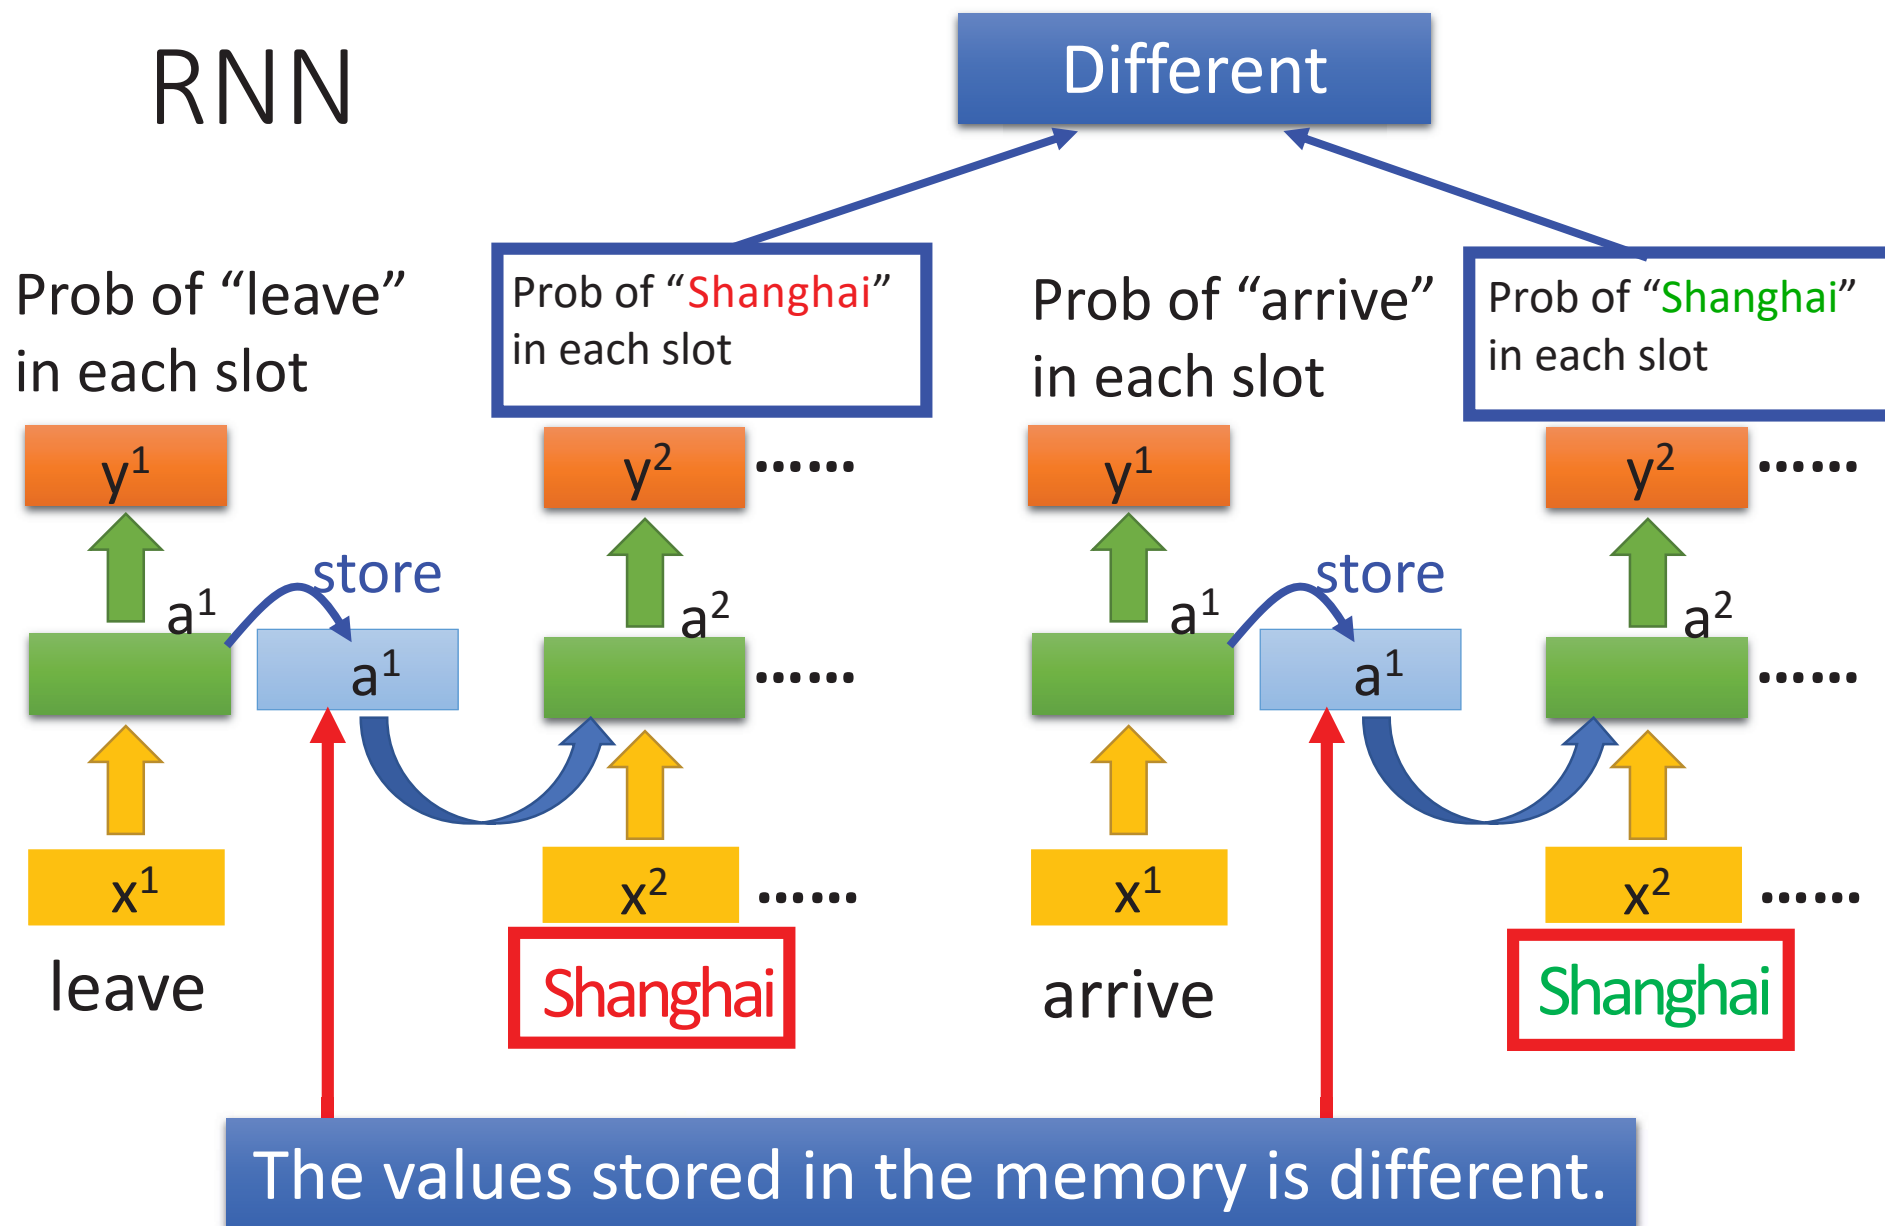
Probability of  
“arrive” in each slot

Probability of  
“**Shanghai**” in each slot

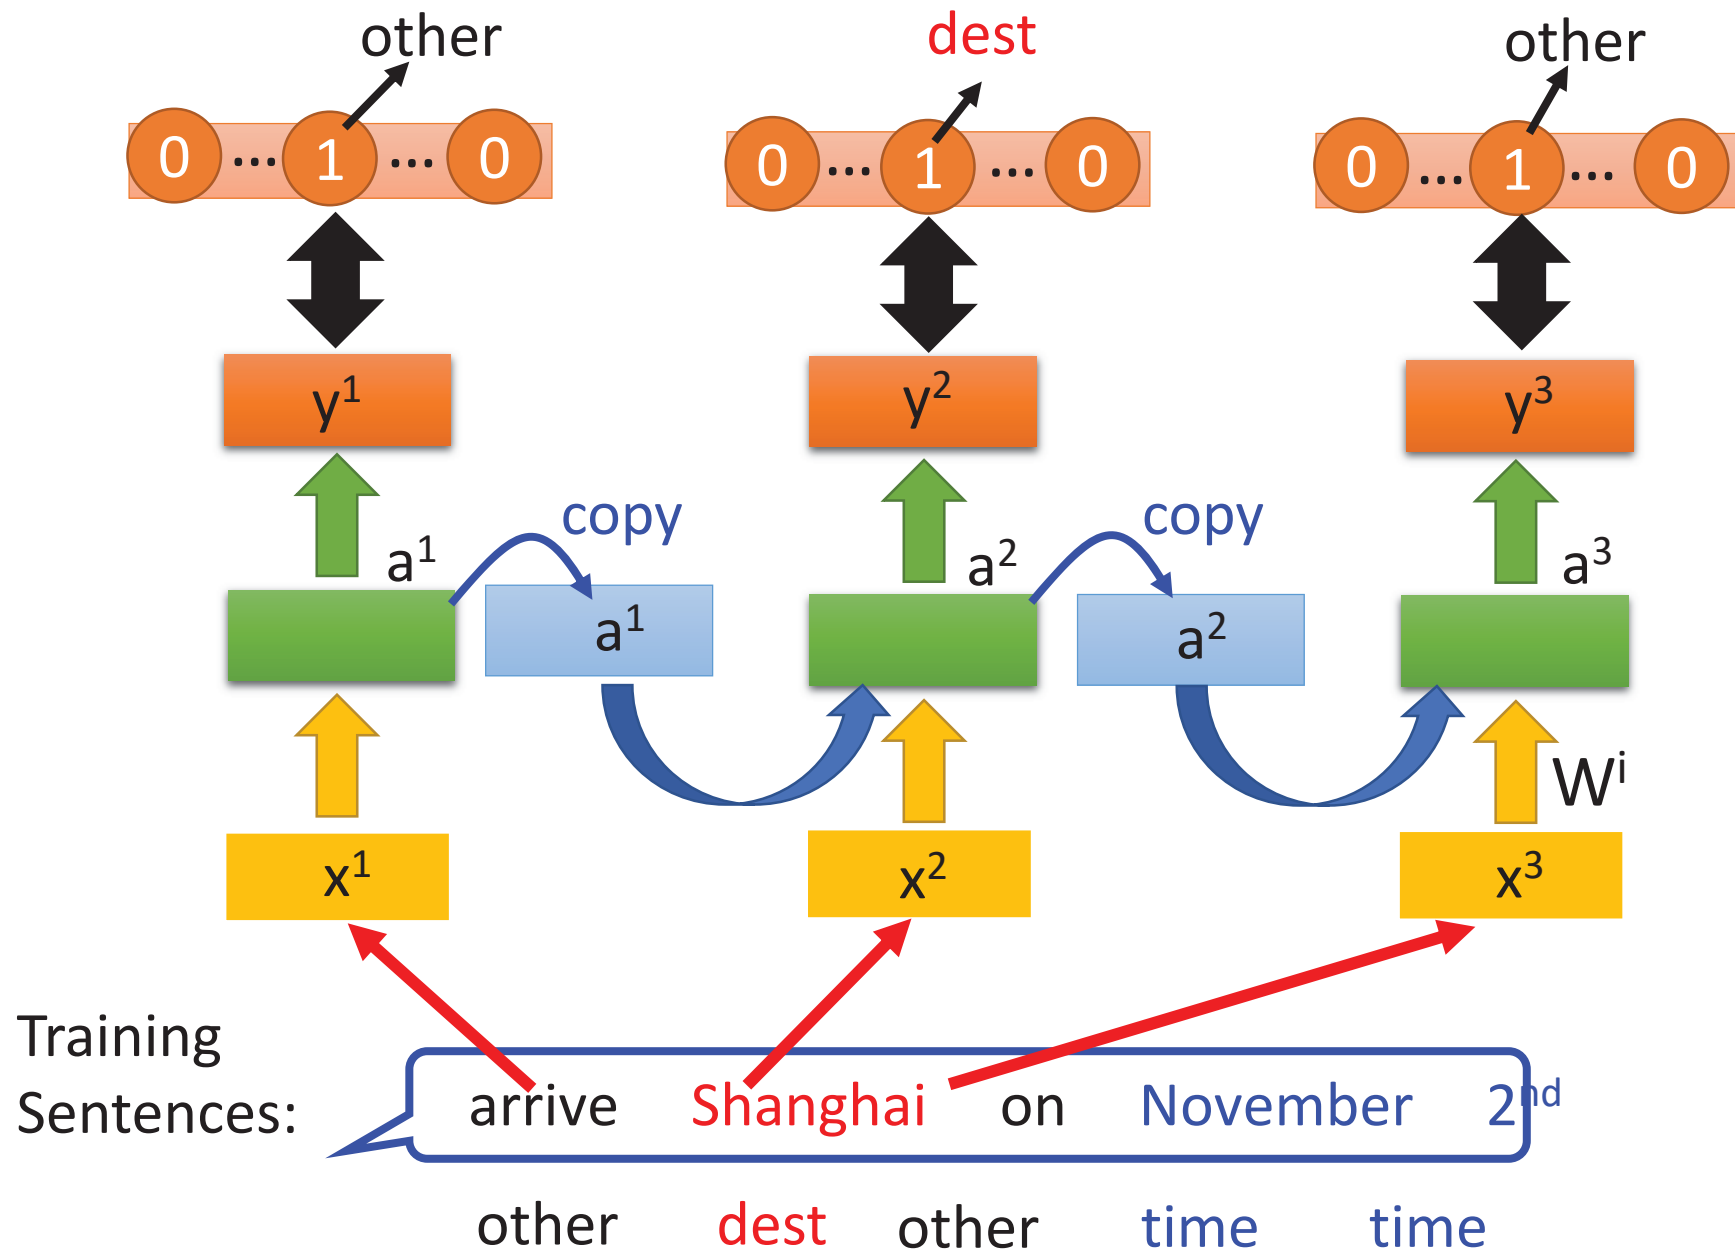
Probability of  
“on” in each slot



# RNN

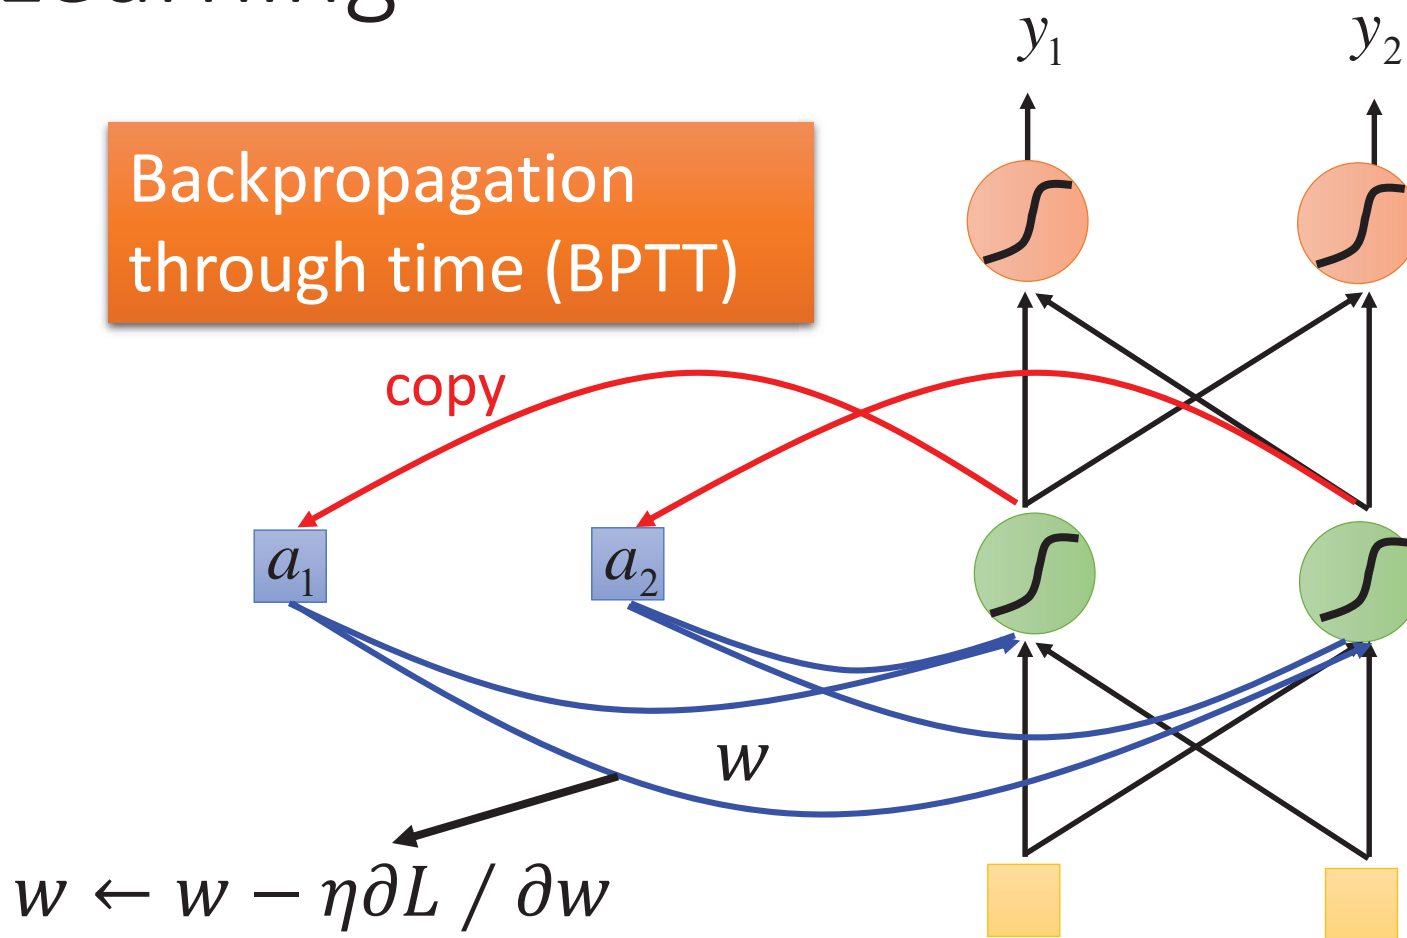


# Learning Target





# Learning



RNN Learning is very difficult in practice.

# vanishing/exploding gradient problem (1)

- Similar but simpler RNN formulation:

$$h_t = W f(h_{t-1}) + W^{(hx)} x_{[t]}$$

$$\hat{y}_t = W^{(S)} f(h_t)$$

- Total error is the sum of each error at time steps  $t$

$$\frac{\partial E}{\partial W} = \sum_{t=1}^T \frac{\partial E_t}{\partial W}$$

- Hardcore chain rule application:

$$\frac{\partial E_t}{\partial W} = \sum_{k=1}^t \frac{\partial E_t}{\partial y_t} \frac{\partial y_t}{\partial h_t} \frac{\partial h_t}{\partial h_k} \frac{\partial h_k}{\partial W}$$

# vanishing/exploding gradient problem (2)

- Similar to backprop but less efficient formulation
- Useful for analysis, we'll look at:

$$\frac{\partial E_t}{\partial W} = \sum_{k=1}^t \frac{\partial E_t}{\partial y_t} \frac{\partial y_t}{\partial h_t} \boxed{\frac{\partial h_t}{\partial h_k}} \frac{\partial h_k}{\partial W}$$

- Remember:  $h_t = W f(h_{t-1}) + W^{(hx)} x_{[t]}$
- More chain rule, remember:


$$\frac{\partial h_t}{\partial h_k} = \prod_{j=k+1}^t \frac{\partial h_j}{\partial h_{j-1}}$$

- Each partial is a Jacobian:

$$\frac{d\mathbf{f}}{d\mathbf{x}} = \begin{bmatrix} \frac{\partial \mathbf{f}}{\partial x_1} & \cdots & \frac{\partial \mathbf{f}}{\partial x_n} \end{bmatrix} = \begin{bmatrix} \frac{\partial f_1}{\partial x_1} & \cdots & \frac{\partial f_1}{\partial x_n} \\ \vdots & \ddots & \vdots \\ \frac{\partial f_m}{\partial x_1} & \cdots & \frac{\partial f_m}{\partial x_n} \end{bmatrix}$$

# vanishing/exploding gradient problem (3)

- From previous slide:  $\frac{\partial h_t}{\partial h_k} = \prod_{j=k+1}^t \frac{\partial h_j}{\partial h_{j-1}}$
- Remember:  $h_t = W f(h_{t-1}) + W^{(hx)} x_{[t]}$



- To compute Jacobian, derive each element of matrix:  $\frac{\partial h_{j,m}}{\partial h_{j-1,n}}$

$$\frac{\partial h_t}{\partial h_k} = \prod_{j=k+1}^t \frac{\partial h_j}{\partial h_{j-1}} = \prod_{j=k+1}^t W^T \text{diag}[f'(h_{j-1})]$$

- Where:  $\text{diag}(z) = \begin{pmatrix} z_1 & & & \\ & z_2 & & 0 \\ & & \ddots & \\ & 0 & & z_{n-1} & \\ & & & & z_n \end{pmatrix}$

Check at home  
that you understand  
the diag matrix  
formulation

# vanishing/exploding gradient problem (4)

- Analyzing the norms of the Jacobians, yields:

$$\left\| \frac{\partial h_j}{\partial h_{j-1}} \right\| \leq \|W^T\| \|\text{diag}[f'(h_{j-1})]\| \leq \beta_W \beta_h$$

- Where we defined  $\beta$ 's as upper bounds of the norms
- The gradient is a product of Jacobian matrices, each associated with a step in the forward computation.

$$\left\| \frac{\partial h_t}{\partial h_k} \right\| = \left\| \prod_{j=k+1}^t \frac{\partial h_j}{\partial h_{j-1}} \right\| \leq (\beta_W \beta_h)^{t-k}$$


- This can become very small or very large quickly [Bengio et al 1994], and the locality assumption of gradient descent breaks down. → **Vanishing or exploding gradient**

# Solve vanishing/exploding gradient

- Clip gradients to a maximum value



```
 $\hat{\mathbf{g}} \leftarrow \frac{\partial \mathcal{E}}{\partial \theta}$   
if  $\|\hat{\mathbf{g}}\| \geq \text{threshold}$  then  
   $\hat{\mathbf{g}} \leftarrow \frac{\text{threshold}}{\|\hat{\mathbf{g}}\|} \hat{\mathbf{g}}$   
end if
```


- Truncated gradient
  - Only use recent information
- Relu or Softplus  
 $\text{Softplus}(x) = \log(1 + e^x)$   $f(x) = \max(0, x)$
- Gate: GRU or LSTM



# Long Short Term Memory (LSTM)

- Add hidden cell as memory C
- Inertia |  $\mathbf{c}_t = \mathbf{i}_t \odot \mathbf{u}_t + \mathbf{f}_t \odot \mathbf{c}_{t-1}$
- Avoid gradient vanishing (gradient derivation)





$$\mathbf{i}_t = \sigma(\mathbf{W}^{(i)}\mathbf{x}_t + \mathbf{U}^{(i)}\mathbf{h}_{t-1} + \mathbf{b}^{(i)})$$

$$\mathbf{o}_t = \sigma(\mathbf{W}^{(o)}\mathbf{x}_t + \mathbf{U}^{(o)}\mathbf{h}_{t-1} + \mathbf{b}^{(o)}) \quad \frac{\partial \mathbf{c}_t}{\partial \mathbf{c}_{t-1}} = \sigma(\mathbf{W}^{(f)}\mathbf{x}_t + \mathbf{U}^{(f)}\mathbf{h}_{t-1} + \mathbf{b}^{(f)})$$

$$\mathbf{f}_t = \sigma(\mathbf{W}^{(f)}\mathbf{x}_t + \mathbf{U}^{(f)}\mathbf{h}_{t-1} + \mathbf{b}^{(f)})$$

$$\mathbf{u}_t = \tanh(\mathbf{W}^{(u)}\mathbf{x}_t + \mathbf{U}^{(u)}\mathbf{h}_{t-1} + \mathbf{b}^{(u)})$$


$$\mathbf{c}_t = \mathbf{i}_t \odot \mathbf{u}_t + \mathbf{f}_t \odot \mathbf{c}_{t-1}$$

$$\mathbf{h}_t = \mathbf{o}_t \odot \tanh(\mathbf{c})$$



# peephole connections

- Gates are related to memories



$$f_t = \sigma(W_f \cdot [C_{t-1}, h_{t-1}, x_t] + b_f)$$

$$i_t = \sigma(W_i \cdot [C_{t-1}, h_{t-1}, x_t] + b_i)$$

$$o_t = \sigma(W_o \cdot [C_t, h_{t-1}, x_t] + b_o)$$


$$\mathbf{u}_t = \tanh(\mathbf{W}^{(u)} \mathbf{x}_t + \mathbf{U}^{(u)} \mathbf{h}_{t-1} + \mathbf{b}^{(u)})$$

$$\mathbf{c}_t = \mathbf{i}_t \odot \mathbf{u}_t + \mathbf{f}_t \odot \mathbf{c}_{t-1}$$

$$\mathbf{h}_t = \mathbf{o}_t \odot \tanh(\mathbf{c})$$

# Gated Recurrent Unit (GRU)

- Combine forget gate and input gate as update gate  $z_t$



$$z_t = \sigma(W_z \cdot [h_{t-1}, x_t])$$

$$r_t = \sigma(W_r \cdot [h_{t-1}, x_t])$$

$$\tilde{h}_t = \tanh(W \cdot [r_t * h_{t-1}, x_t])$$

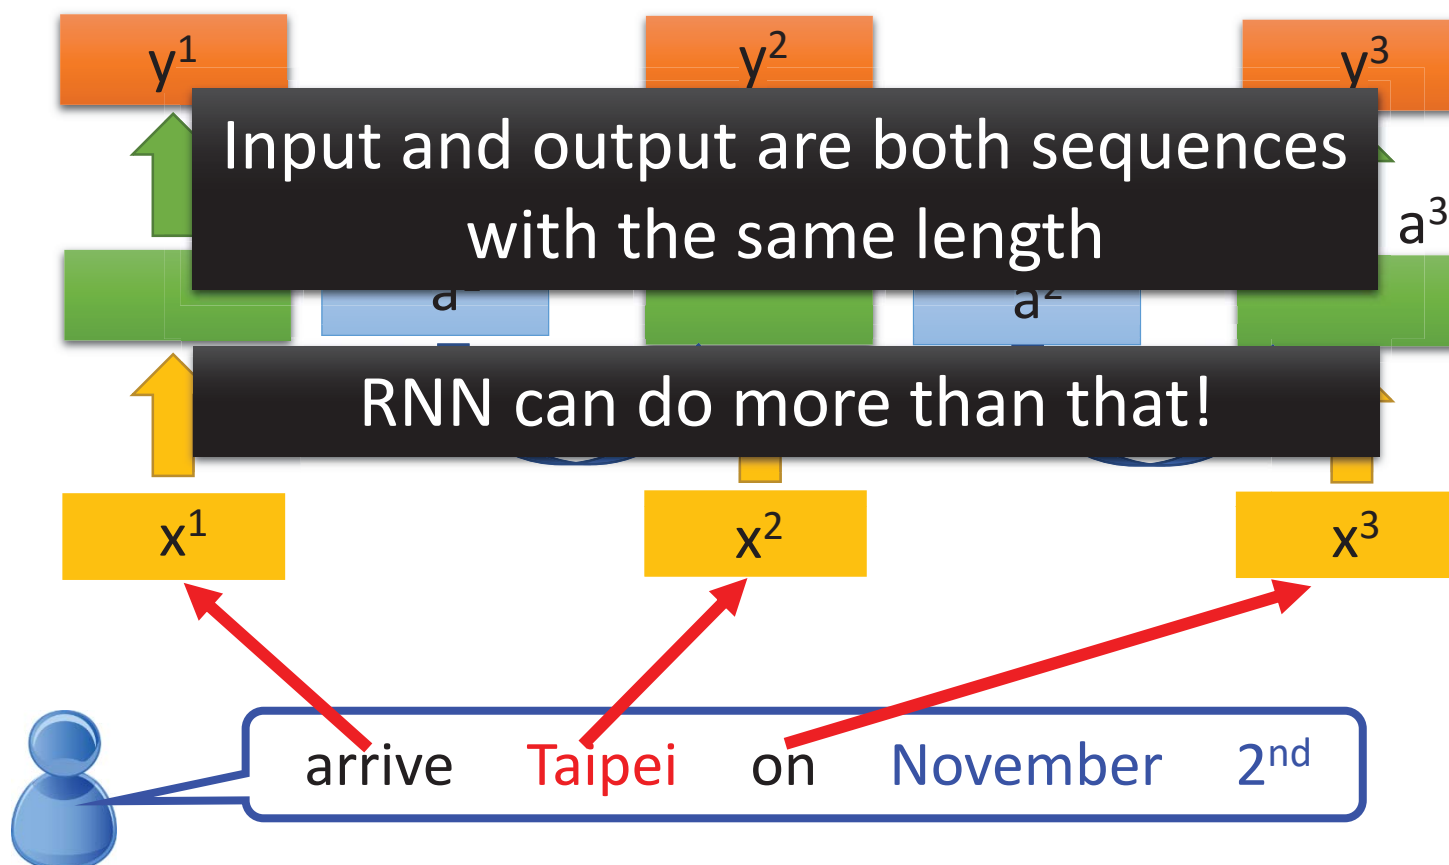
$$h_t = (1 - z_t) * h_{t-1} + z_t * \tilde{h}_t$$

# More Applications .....

Probability of  
“arrive” in each slot

Probability of  
“**Taipei**” in each slot

Probability of  
“on” in each slot




# Many to one

Keras Example:

[https://github.com/fchollet/keras/blob/master/examples/imdb\\_lstm.py](https://github.com/fchollet/keras/blob/master/examples/imdb_lstm.py)

- Input is a vector sequence, but output is only one vector

## Sentiment Analysis



# Many to Many (Output is shorter)

- Both input and output are both sequences, **but the output is shorter.**
  - E.g. **Speech Recognition**

Problem?

Why can't it be  
“好棒棒”


Output: “好棒” (character sequence)



Trimming

好 好 好 棒 棒 棒 棒 棒

Input:




(vector  
sequence)



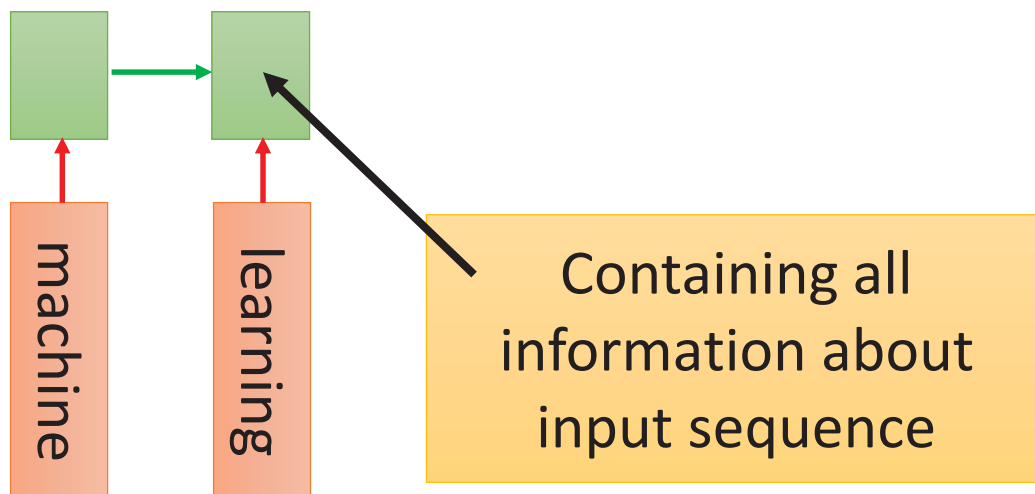
# Many to Many (Output is shorter)

- Both input and output are both sequences, **but the output is shorter.**
- Connectionist Temporal Classification (CTC) [Alex Graves, ICML'06][Alex Graves, ICML'14][Haşim Sak, Interspeech'15][Jie Li, Interspeech'15][Andrew Senior, ASRU'15]




# Many to Many (No Limitation)

- Both input and output are both sequences **with different lengths.** → **Sequence to sequence learning**
  - E.g. **Machine Translation** (machine learning → 機器學習)



# Many to Many (No Limitation)

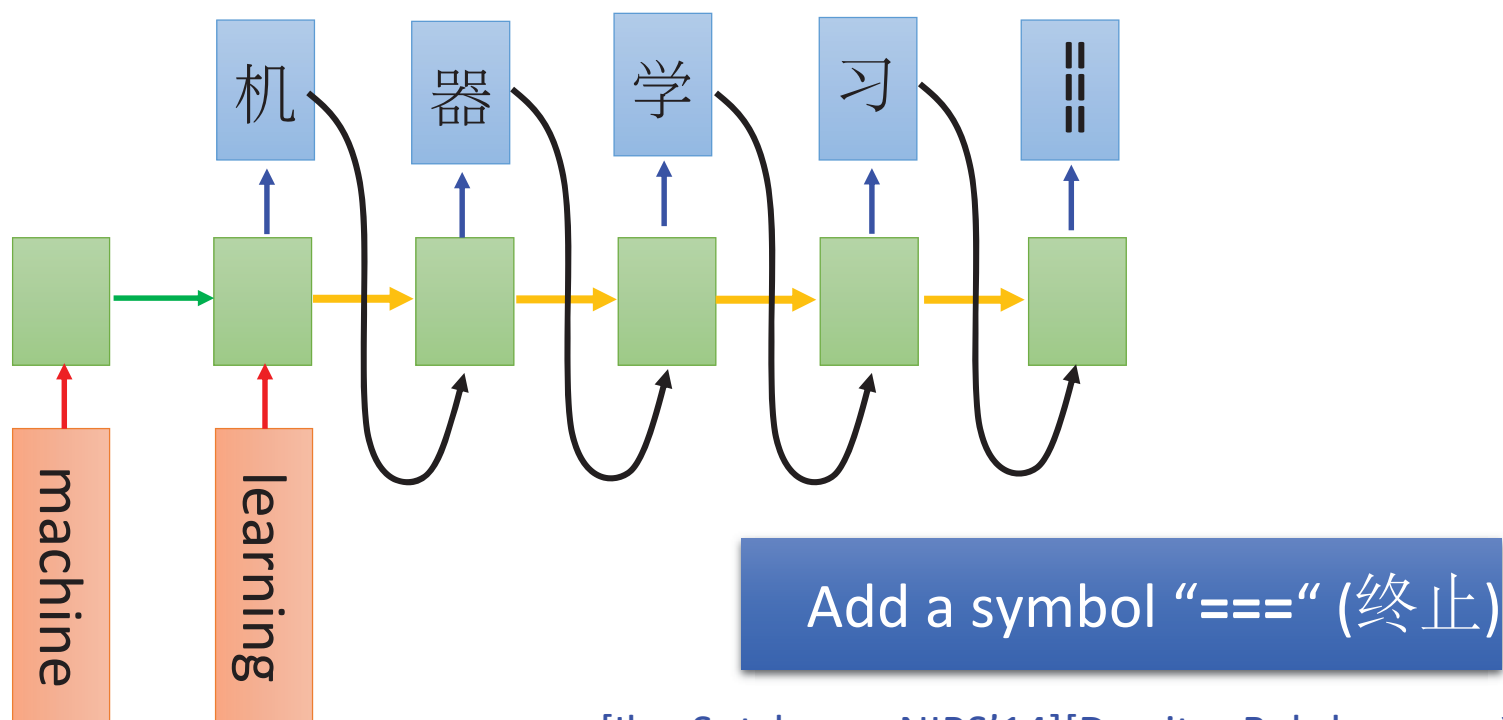
- Both input and output are both sequences with different lengths. → Sequence to sequence learning
  - E.g. Machine Translation (machine learning → 机器学习)





# Many to Many (No Limitation)

- Both input and output are both sequences **with different lengths**. → **Sequence to sequence learning**
  - E.g. **Machine Translation** (machine learning → 机器学习)



[Ilya Sutskever, NIPS'14][Dzmitry Bahdanau, arXiv'15]

# One to Many

- Input an image, but output a sequence of words

[Kelvin Xu, arXiv'15][Li Yao, ICCV'15]

